



UL HPC School 2017

PS5: Advanced Scheduling with SLURM and OAR on UL HPC clusters

UL High Performance Computing (HPC) Team

V. Plugaru

University of Luxembourg (UL), Luxembourg

<http://hpc.uni.lu>

Latest versions available on Github:



UL HPC tutorials:

<https://github.com/ULHPC/tutorials>

UL HPC School:

<http://hpc.uni.lu/hpc-school/>

PS5 tutorial sources:

https://github.com/ULHPC/tutorials/tree/devel/advanced/advanced_scheduling





Summary

- 1 **Introduction**
- 2 SLURM workload manager
SLURM concepts and design for iris
Running jobs with SLURM
- 3 OAR and SLURM
- 4 Conclusion

Main Objectives of this Session



- **Design and usage of SLURM**

- ↪ cluster workload manager of the UL HPC iris cluster
- ↪ ... and future HPC systems

The tutorial will show you:

- the way SLURM was **configured**, **accounting** and **permissions**
- **common** and **advanced** SLURM tools and commands
 - ↪ srun, sbatch, squeue etc.
 - ↪ job specification
 - ↪ SLURM job types
 - ↪ comparison of SLURM (iris) and OAR (gaia & chaos)
- SLURM generic **launchers** you can use for your own jobs

Documentation & comparison to OAR

<https://hpc.uni.lu/users/docs/scheduler.html>



Summary

- 1 Introduction
- 2 SLURM workload manager**
SLURM concepts and design for iris
Running jobs with SLURM
- 3 OAR and SLURM
- 4 Conclusion

SLURM - core concepts

- SLURM manages user jobs with the following **key characteristics**:

→ set of **requested resources**:

- ✓ number of computing resources: **nodes** (including all their CPUs and cores) or **CPUs** (including all their cores) or **cores**
- ✓ amount of **memory**: either per node or per (logical) CPU
- ✓ **(wall)time** needed for the user's tasks to complete their work

→ a requested node **partition** (job queue)

→ a requested **quality of service** (QoS) level which grants users specific accesses

→ a requested **account** for accounting purposes

- Example**: run an interactive job

Alias: `si [...]`

```
(access)$ srun -p interactive --qos qos-interactive --pty bash
(node)$ echo $SLURM_JOBID
2058
```

Simple interactive job running under SLURM

SLURM - job example (I)

```
$ scontrol show job 2058
JobId=2058 JobName=bash
  UserId=vplugaru(5143) GroupId=clusterusers(666) MCS_label=N/A
  Priority =100 Nice=0 Account=ulhpc QOS=qos-interactive
5  JobState=RUNNING Reason=None Dependency=(null)
  Requeue=1 Restarts=0 BatchFlag=0 Reboot=0 ExitCode=0:0
  RunTime=00:00:08 TimeLimit=00:05:00 TimeMin=N/A
  SubmitTime=2017-06-09T16:49:42 EligibleTime=2017-06-09T16:49:42
10 StartTime=2017-06-09T16:49:42 EndTime=2017-06-09T16:54:42 Deadline=N/A
  PreemptTime=None SuspendTime=None SecsPreSuspend=0
  Partition = interactive AllocNode:Sid=access2:163067
  ReqNodeList=(null) ExcNodeList=(null)
  NodeList=iris-081
  BatchHost=iris-081
15 NumNodes=1 NumCPUs=1 NumTasks=1 CPUs/Task=1 ReqB:S:C:T=0:0:*:*
  TRES=cpu=1,mem=4G,node=1
  Socks/Node=* NtasksPerN:B:S:C=1:0:*:* CoreSpec=*
  MinCPUsNode=1 MinMemoryCPU=4G MinTmpDiskNode=0
  Features=(null) DelayBoot=00:00:00
20 Gres=(null) Reservation=(null)
  OverSubscribe=OK Contiguous=0 Licenses=(null) Network=(null)
  Command=bash
  WorkDir=/mnt/irisgpf/users/vplugaru
  Power=
```

Simple interactive job running under SLURM

SLURM - job example (II)

- Many metrics available during and after job execution
 - including energy (J) – but with caveats
 - job **steps** counted individually
 - enabling advanced application debugging and optimization
- Job information available in easily parseable format (add -p/-P)

```

$ sacct -j 2058 --format=account,user,jobid,jobname,partition,state
  Account      User      JobID  JobName  Partition  State
    ulhpc    vplugaru      2058      bash  interacti +  COMPLETED

5  $ sacct -j 2058 --format=elapsed,elapsedraw,start,end
    Elapsed ElapsedRaw      Start      End
    00:02:56      176 2017-06-09T16:49:42 2017-06-09T16:52:38

10 $ sacct -j 2058 --format=maxrss,maxvmsize,consumedenergy,consumedenergyraw,nnodes,ncpus,nodelist
    MaxRSS MaxVMSize ConsumedEnergy ConsumedEnergyRaw NNodes NCPUS  NodeList
        0   299660K      17.89K      17885.000000        1        1    iris -081
  
```

Job metrics after execution ended

SLURM - design for iris (I)

Partition	# Nodes	Default time	Max time	Max nodes/user
batch*	88 (82%)	0-2:0:0	5-0:0:0	unlimited
interactive	10 (9%)	0-1:0:0	0-4:0:0	2
long	10 (9%)	0-2:0:0	30-0:0:0	2

SLURM - design for iris (I)

Partition	# Nodes	Default time	Max time	Max nodes/user
batch*	88 (82%)	0-2:0:0	5-0:0:0	unlimited
interactive	10 (9%)	0-1:0:0	0-4:0:0	2
long	10 (9%)	0-2:0:0	30-0:0:0	2

QoS	User group	Max cores	Max jobs/user
qos-besteffort	ALL	no limit	
qos-batch	ALL	1064	100
qos-interactive	ALL	224	10
qos-long	ALL	224	10
qos-batch-001	private	1400	100
qos-interactive-001	private	56	10
qos-long-001	private	56	10

SLURM - design for iris (II)

- **Default partition: `batch`**, meant to receive most user jobs
 - ↪ we hope to see majority of user jobs being able to scale
- All partitions have a correspondingly named **QOS**
 - ↪ granting resource access (**long** – qos-long)
 - ↪ any job is tied to one QOS (user specified or inferred)
 - ↪ automation in place to select QOS based on partition

SLURM - design for iris (II)

- **Default partition:** **batch**, meant to receive most user jobs
 - ↪ we hope to see majority of user jobs being able to scale
- All partitions have a correspondingly named **QOS**
 - ↪ granting resource access (**long** – **qos-long**)
 - ↪ any job is tied to one QOS (user specified or inferred)
 - ↪ automation in place to select QOS based on partition
- Preemptible **besteffort** QOS available for **batch** and **interactive** partitions (but not for **long**)
 - ↪ meant to ensure maximum resource utilization
 - ↪ should be used together with checkpointable software
- QOSs specific to particular group accounts exist (discussed later)
 - ↪ granting additional accesses to platform contributors

SLURM - design for iris (III)

- **Backfill** scheduling for efficiency
 - ↪ **multifactor job priority** (size, age, fairshare, QOS, ...)
 - ↪ currently weights set for: job age, partition and fair-share
 - ↪ other factors/decay to be tuned **after observation** period
 - ✓ with more user jobs in the queues
- Resource selection: **consumable resources**
 - ↪ **cores and memory** as consumable (per-core scheduling)
 - ↪ block distribution for cores (best-fit algorithm)
 - ↪ default memory/core: 4GB (4.1GB maximum, rest is for OS)

SLURM - design for iris (III)

- **Backfill** scheduling for efficiency
 - ↪ **multifactor job priority** (size, age, fairshare, QOS, ...)
 - ↪ currently weights set for: job age, partition and fair-share
 - ↪ other factors/decay to be tuned **after observation** period
 - ✓ with more user jobs in the queues
- Resource selection: **consumable resources**
 - ↪ **cores and memory** as consumable (per-core scheduling)
 - ↪ block distribution for cores (best-fit algorithm)
 - ↪ default memory/core: 4GB (4.1GB maximum, rest is for OS)
- Reliable user process tracking with **cgroups**
 - ↪ cpusets used to constrain cores and RAM (no swap allowed)
 - ↪ task affinity used to bind tasks to cores (hwloc based)
- Hierarchical tree topology defined (for the network)
 - ↪ for optimized job resource allocation

SLURM - design for iris (III)

- **Backfill** scheduling for efficiency
 - **multifactor job priority** (size, age, fairness, share)
 - currently weights set for: job age, job size, job priority, job share
 - other factors/decay to be tuned for specific use case and period
 - ✓ with more user jobs in the queue
- Resource selection: **consumption**
 - **cores and memory** as primary resource (for core scheduling)
 - block distribution for memory (first fit algorithm)
 - default memory/cores ratio (user can set maximum, rest is for OS)
- Reliable user process management with **cgroups**
 - cpuset for CPU and memory (no swap allowed)
 - task placement and tasks to cores (hwloc based)
- Hierarchical topology defined (for the network)
 - for efficient job resource allocation

**Help will be needed on your part
to optimize your job parameters!**

A note on job priority

```
Job_priority =  
    (PriorityWeightAge) * (age_factor) +  
    (PriorityWeightFairshare) * (fair-share_factor) +  
    (PriorityWeightJobSize) * (job_size_factor) +  
    (PriorityWeightPartition) * (partition_factor) +  
    (PriorityWeightQOS) * (QOS_factor) +  
    SUM(TRES_weight_cpu * TRES_factor_cpu,  
        TRES_weight_<type> * TRES_factor_<type>,  
        ...)
```

- TRES – Trackable RESources
 - ↪ CPU, Energy, Memory and Node tracked by default All details at slurm.schedmd.com/priority_multifactor.html
- The corresponding weights and reset periods we **need to tune**
 - ↪ we require (your!) real application usage to optimize them

SLURM - design for iris (IV)

Some details on job permissions...

- Partition limits + association-based rule enforcement
 - ↳ association settings in SLURM's accounting database
- **QOS** limits imposed, e.g. you will see (QOSGrpCpuLimit)
- Only users with existing **associations** able to run jobs

SLURM - design for iris (IV)

Some details on job permissions...

- Partition limits + association-based rule enforcement
 - ↪ association settings in SLURM's accounting database
 - **QOS** limits imposed, e.g. you will see (QOSGrpCpuLimit)
 - Only users with existing **associations** able to run jobs
-
- **Best-effort** jobs possible through preemptible QOS: **qos-besteffort**
 - ↪ of lower priority and preemptible by all other QOS
 - ↪ preemption mode is **requeue**, requeueing enabled by default

SLURM - design for iris (IV)

Some details on job permissions...

- Partition limits + association-based rule enforcement
 - ↪ association settings in SLURM's accounting database
 - **QOS** limits imposed, e.g. you will see (QOSGrpCpuLimit)
 - Only users with existing **associations** able to run jobs
-
- **Best-effort** jobs possible through preemptible QOS: **qos-besteffort**
 - ↪ of lower priority and preemptible by all other QOS
 - ↪ preemption mode is **requeue**, requeueing enabled by default
 - **On metrics**: Accounting & profiling data for jobs sampled every 30s
 - ↪ tracked: cpu, mem, energy
 - ↪ energy data retrieved through the **RAPL mechanism**
 - ✓ **caveat**: for energy not all hw. that may consume power is monitored with RAPL (CPUs, GPUs and DRAM are included)

SLURM - design for iris (V)

- **On tightly coupled parallel jobs (MPI)**

- ↳ Process Management Interface (PMI 2) highly recommended
- ↳ PMI2 used for better scalability and performance
 - ✓ faster application launches
 - ✓ tight integration w. SLURM's job steps mechanism (& metrics)
 - ✓ we are also testing **PMIx** (PMI Exascale) support

SLURM - design for iris (V)

- On tightly coupled parallel jobs (MPI)

- ↪ Process Management Interface (PMI 2) highly recommended
- ↪ PMI2 used for better scalability and performance
 - ✓ faster application launches
 - ✓ tight integration w. SLURM's job steps mechanism (& metrics)
 - ✓ we are also testing **PMIx** (PMI Exascale) support
- ↪ PMI2 enabled in default software set for IntelMPI and OpenMPI
 - ✓ requires minimal adaptation in your workflows
 - ✓ replace **mpirun** with SLURM's **srun** (at minimum)
 - ✓ if you compile/install your own MPI you'll need to configure it
- ↪ **Example:** https://hpc.uni.lu/users/docs/slurm_launchers.html

SLURM - design for iris (V)

- **On tightly coupled parallel jobs (MPI)**

- Process Management Interface (PMI 2) highly recommended
- PMI2 used for better scalability and performance
 - ✓ faster application launches
 - ✓ tight integration w. SLURM's job steps mechanism (& metrics)
 - ✓ we are also testing **PMIx** (PMI Exascale) support
- PMI2 enabled in default software set for IntelMPI and OpenMPI
 - ✓ requires minimal adaptation in your workflows
 - ✓ replace **mpirun** with SLURM's **srun** (at minimum)
 - ✓ if you compile/install your own MPI you'll need to configure it
- **Example:** https://hpc.uni.lu/users/docs/slurm_launchers.html

- **SSH-based connections** between computing nodes still **possible**

- other MPI implementations can still use `ssh` as launcher
 - ✓ but **really shouldn't need to**, PMI2 support is everywhere
- user jobs are tracked, no job == no access to node

SLURM - design for iris (VI)

ULHPC customizations through plugins

- Job submission rule / filter
 - ↳ for now: QOS initialization (if needed)
 - ↳ more rules to come (group credits, node checks, etc.)
- Per-job temporary directories creation & cleanup
 - ↳ better security and privacy, using kernel namespaces and binding
 - ↳ /tmp & /var/tmp are /tmp/\$jobid.\$rstcnt/[tmp,var_tmp]
 - ↳ transparent for apps. ran through srun
 - ↳ **apps. ran with ssh cannot be attached, will see base /tmp!**
- X11 forwarding (GUI applications)
 - ↳ enabled with --x11 parameter to srun/salloc
 - ↳ **currently being rewritten to play nice with per-job tmpdir**
 - ✓ workaround: create job and ssh -X to head node (need to propagate job environment)

SLURM - design for iris (VII)


Software licenses in SLURM

- Allinea Forge and Performance Reports for now
 - ↪ static allocation in SLURM configuration
 - ↪ dynamic checks for FlexNet / RLM based apps. coming later
- Number and utilization state can be checked with:
 - ↪ `scontrol show licenses`
- Use not enforced, **honor system** applied
 - ↪ `srun [...] -L $licname:$licnumber`

```
$> srun -N 1 -n 28 -p interactive -L forge:28 --pty bash -i
```


SLURM - bank (group) accounts

- Hierarchical **bank (group) accounts**
- UL as root account, then underneath accounts for the 3 Faculties and 3 ICs
- All Prof., Group leaders and above have **bank accounts**, linked to a Faculty or IC
 - ↳ with their own name: **Name.Surname**
- All **user accounts** linked to a bank account
 - ↳ including Profs.'s own user
- Iris accounting DB contains over
 - ↳ 75 group accounts from all Faculties/ICs
 - ↳ comprising 477 users



Allows better usage tracking and reporting than was possible before.



SLURM - brief commands overview

- **squeue**: view queued jobs
- **sinfo**: view partition and node info.
- **sbatch**: submit job for batch (scripted) execution
- **srun**: submit interactive job, run (parallel) job step
- **scancel**: cancel queued jobs

SLURM - brief commands overview

- **squeue**: view queued jobs
 - **sinfo**: view partition and node info.
 - **sbatch**: submit job for batch (scripted) execution
 - **srun**: submit interactive job, run (parallel) job step
 - **scancel**: cancel queued jobs
-
- **scontrol**: detailed control and info. on jobs, queues, partitions
 - **sstat**: view system-level utilization (memory, I/O, energy)
 - ↪ for running jobs / job steps
 - **sacct**: view system-level utilization
 - ↪ for completed jobs / job steps (accounting DB)
 - **sacctmgr**: view and manage SLURM accounting data

SLURM - brief commands overview

- **squeue**: view queued jobs
 - **sinfo**: view partition and node info.
 - **sbatch**: submit job for batch (scripted) execution
 - **srun**: submit interactive job, run (parallel) job step
 - **scancel**: cancel queued jobs
-
- **scontrol**: detailed control and info. on jobs, queues, partitions
 - **sstat**: view system-level utilization (memory, I/O, energy)
 - ↪ for running jobs / job steps
 - **sacct**: view system-level utilization
 - ↪ for completed jobs / job steps (accounting DB)
 - **sacctmgr**: view and manage SLURM accounting data
-
- **sprio**: view job priority factors
 - **sshare**: view accounting share info. (usage, fair-share, etc.)

SLURM - basic commands

Action	SLURM command
Submit passive/batch job	<code>sbatch \$script</code>
Start interactive job	<code>srun --pty bash -i</code>
Queue status	<code>squeue</code>
User job status	<code>squeue -u \$user</code>
Specific job status (detailed)	<code>scontrol show job \$jobid</code>
Job metrics (detailed)	<code>sstat --job \$jobid -l</code>
Job accounting status (detailed)	<code>sacct --job \$jobid -l</code>
Delete (running/waiting) job	<code>scancel \$jobid</code>
Hold job	<code>scontrol hold \$jobid</code>
Resume held job	<code>scontrol release \$jobid</code>
Node list and their properties	<code>scontrol show nodes</code>
Partition list, status and limits	<code>sinfo</code>

QOS deduced if not specified, partition needs to be set if not "batch"

SLURM - basic options for sbatch/srun

Action	sbatch/srun option
Request \$n distributed nodes	-N \$n
Request \$m memory per node	--mem=\$mGB
Request \$mc memory per core (logical cpu)	--mem-per-cpu=\$mcGB
Request job walltime	--time=d-hh:mm:ss
Request \$tn tasks per node	--ntasks-per-node=\$tn
Request \$ct cores per task (multithreading)	-c \$ct
Request \$nt total # of tasks	-n \$nt
Request to start job at specific \$time	--begin \$time
Specify job name as \$name	-J \$name
Specify job partition	-p \$partition
Specify QOS	--qos \$qos
Specify account	-A \$account
Specify email address	--mail-user=\$email
Request email on event	--mail-type=all[,begin,end,fail]
Use the above actions in a batch script	#SBATCH \$option

SLURM - basic options for sbatch/srun

Action	sbatch/srun option
Request \$n distributed nodes	-N \$n
Request \$m memory per node	--mem=\$mGB
Request \$mc memory per core (logical cpu)	--mem-per-cpu=\$mcGB
Request job walltime	--time=d-hh:mm:ss
Request \$tn tasks per node	--ntasks-per-node=\$tn
Request \$ct cores per task (multithreading)	-c \$ct
Request \$nt total # of tasks	-n \$nt
Request to start job at specific \$time	--begin \$time
Specify job name as \$name	-J \$name
Specify job partition	-p \$partition
Specify QOS	--qos \$qos
Specify account	-A \$account
Specify email address	--mail-user=\$email
Request email on event	--mail-type=all[,begin,end,fail]
Use the above actions in a batch script	#SBATCH \$option

- Diff. between **-N**, **-c**, **-n**, **--ntasks-per-node**, **--ntasks-per-core** ?
- Normally you'd specify **-N** and **--ntasks-per-node**
 - ↪ fix the latter to 1 and add **-c** for MPI+OpenMP jobs
- If your application is scalable, just **-n** might be enough
 - ↪ iris is homogeneous (for now)

SLURM - more options for sbatch/srun

Start job when... (dependencies)	sbatch/srun option
these other jobs have started	-d after:\$jobid1:\$jobid2
these other jobs have ended	-d afterany:\$jobid1:\$jobid2
these other jobs have ended with no errors	-d afterok:\$jobid1:\$jobid2
these other jobs have ended with errors	-d afternok:\$jobid1:\$jobid2
all other jobs with the same name have ended	-d singleton

Job dependencies and especially "singleton" can be very useful!

SLURM - more options for sbatch/srun

Start job when... (dependencies)	sbatch/srun option
these other jobs have started	-d after:\$jobid1:\$jobid2
these other jobs have ended	-d afterany:\$jobid1:\$jobid2
these other jobs have ended with no errors	-d afterok:\$jobid1:\$jobid2
these other jobs have ended with errors	-d afternok:\$jobid1:\$jobid2
all other jobs with the same name have ended	-d singleton

Job dependencies and especially "singleton" can be very useful!

Allocate job at... (specified time)	sbatch/srun option
exact time today	--begin=16:00
tomorrow	--begin=tomorrow
specific time relative to now	--begin=now+2hours
given date and time	--begin=2017-06-23T07:30:00

Jobs run like this will wait as PD – Pending with "(BeginTime)" reason

SLURM - more options for sbatch/srun

Start job when... (dependencies)	sbatch/srun option
these other jobs have started	-d after:\$jobid1:\$jobid2
these other jobs have ended	-d afterany:\$jobid1:\$jobid2
these other jobs have ended with no errors	-d afterok:\$jobid1:\$jobid2
these other jobs have ended with errors	-d afternok:\$jobid1:\$jobid2
all other jobs with the same name have ended	-d singleton

Job dependencies and especially "singleton" can be very useful!

Allocate job at... (specified time)	sbatch/srun option
exact time today	--begin=16:00
tomorrow	--begin=tomorrow
specific time relative to now	--begin=now+2hours
given date and time	--begin=2017-06-23T07:30:00

Jobs run like this will wait as PD – Pending with "(BeginTime)" reason

Other scheduling request	sbatch/srun option
Ask for minimum/maximum # of nodes	-N minnodes-maxnodes
Ask for minimum run time (start job faster)	--time-min=d-hh:mm:ss
Ask to remove job if deadline can't be met	--deadline=YYYY-MM-DD[THH:MM[:SS]]
Run job within pre-created (admin) reservation	--reservation=\$reservationname
Allocate resources as specified job	--jobid=\$jobid

Can use --jobid to connect to running job (different than sattach!)

SLURM - environment variables

- 53 input env. vars. can be used to define job parameters
 - ↪ almost all have a command line equivalent
- up to 59 output env. vars. available within job environment
 - ↪ some common ones:

Description	Environment variable
Job ID	<code>\$SLURM_JOBID</code>
Job name	<code>\$SLURM_JOB_NAME</code>
Name of account under which job runs	<code>\$SLURM_JOB_ACCOUNT</code>
Name of partition job is running in	<code>\$SLURM_JOB_PARTITION</code>
Name of QOS the job is running with	<code>\$SLURM_JOB_QOS</code>
Name of job's advance reservation	<code>\$SLURM_JOB_RESERVATION</code>
Job submission directory	<code>\$SLURM_SUBMIT_DIR</code>
Number of nodes assigned to the job	<code>\$SLURM_NNODES</code>
Name of nodes assigned to the job	<code>\$SLURM_JOB_NODELIST</code>
Number of tasks for the job	<code>\$SLURM_NTASKS</code> or <code>\$SLURM_NPROCS</code>
Number of cores for the job on current node	<code>\$SLURM_JOB_CPUS_PER_NODE</code>
Memory allocated to the job per node	<code>\$SLURM_MEM_PER_NODE</code>
Memory allocated per core	<code>\$SLURM_MEM_PER_CPU</code>
Task count within a job array	<code>\$SLURM_ARRAY_TASK_COUNT</code>
Task ID assigned within a job array	<code>\$SLURM_ARRAY_TASK_ID</code>

Outputting these variables to the job log is essential for bookkeeping!



Usage examples (I)

> Interactive jobs

```
srunk -p interactive --qos qos-interactive --time=0:30 -N2 --ntasks-per-node=4 --pty bash -i  
srunk -p interactive --qos qos-interactive --pty --x11 bash -i  
srunk -p interactive --qos qos-besteffort --pty bash -i
```

Usage examples (I)

> Interactive jobs

```
srunk -p interactive --qos qos-interactive --time=0:30 -N2 --ntasks-per-node=4 --pty bash -i
srunk -p interactive --qos qos-interactive --pty --x11 bash -i
srunk -p interactive --qos qos-besteffort --pty bash -i
```

> Batch jobs

```
sbatch job.sh
sbatch -N 2 job.sh
sbatch -p batch --qos qos-batch job.sh
sbatch -p long --qos qos-long job.sh
sbatch --begin=2017-06-23T07:30:00 job.sh
sbatch -p batch --qos qos-besteffort job.sh
```

Usage examples (I)

> Interactive jobs

```
srun -p interactive --qos qos-interactive --time=0:30 -N2 --ntasks-per-node=4 --pty bash -i
srun -p interactive --qos qos-interactive --pty --x11 bash -i
srun -p interactive --qos qos-besteffort --pty bash -i
```

> Batch jobs

```
sbatch job.sh
sbatch -N 2 job.sh
sbatch -p batch --qos qos-batch job.sh
sbatch -p long --qos qos-long job.sh
sbatch --begin=2017-06-23T07:30:00 job.sh
sbatch -p batch --qos qos-besteffort job.sh
```

Status and details for partitions, nodes, reservations

```
squeue / squeue -l / squeue -la / squeue -l -p batch / squeue -t PD
scontrol show nodes / scontrol show nodes $nodename
sinfo / sinfo -s / sinfo -N
sinfo -T
```

Usage examples (II)

Collecting job information, priority, expected start time

```
scontrol show job $jobid # this is only available while job is in the queue + 5 minutes  
sprio -l  
squeue --start -u $USER
```

Usage examples (II)

Collecting job information, priority, expected start time

```
scontrol show job $jobid # this is only available while job is in the queue + 5 minutes  
sprio -l  
squeue --start -u $USER
```

Running job metrics – sstat tool

```
sstat -j $jobid / sstat -j $jobid -l  
sstat -j $jobid1 --format=AveCPU,AveRSS,AveVMSize,MaxRSS,MaxVMSize  
sstat -p -j $jobid1,$jobid2 --format=AveCPU,AveRSS,AveVMSize,MaxRSS,MaxVMSize
```


Usage examples (II)

Collecting job information, priority, expected start time

```
scontrol show job $jobid # this is only available while job is in the queue + 5 minutes  
sprio -l  
squeue --start -u $USER
```

Running job metrics – sstat tool

```
sstat -j $jobid / sstat -j $jobid -l  
sstat -j $jobid1 --format=AveCPU,AveRSS,AveVMSize,MaxRSS,MaxVMSize  
sstat -p -j $jobid1,$jobid2 --format=AveCPU,AveRSS,AveVMSize,MaxRSS,MaxVMSize
```

Completed job metrics – sacct tool

```
sacct -j $jobid / sacct -j $jobid -l  
sacct -p -j $jobid --format=account,user,jobid,jobname,partition,state,elapsed,elapseddraw,  
\ start,end,maxrss,maxvmsize,consumedenergy,consumedenergyraw,nnodes,ncpus,nodelist  
sacct --starttime 2017-06-12 -u $USER
```

Usage examples (III)

Controlling queued and running jobs

```
scontrol hold $jobid  
scontrol release $jobid  
scontrol suspend $jobid  
scontrol resume $jobid  
scancel $jobid  
scancel -n $jobname  
scancel -u $USER  
scancel -u $USER -p batch  
scontrol requeue $jobid
```

Usage examples (III)

Controlling queued and running jobs

```
scontrol hold $jobid  
scontrol release $jobid  
scontrol suspend $jobid  
scontrol resume $jobid  
scancel $jobid  
scancel -n $jobname  
scancel -u $USER  
scancel -u $USER -p batch  
scontrol requeue $jobid
```

Checking accounting links and QOS available for you

```
sacctmgr show user $USER format=user%20s,defaultaccount%30s  
sacctmgr list association where  
users=$USER format=account%30s,user%20s,qos%120s
```

Usage examples (III)

Controlling queued and running jobs

```
scontrol hold $jobid
scontrol release $jobid
scontrol suspend $jobid
scontrol resume $jobid
scancel $jobid
scancel -n $jobname
scancel -u $USER
scancel -u $USER -p batch
scontrol requeue $jobid
```

Checking accounting links and QOS available for you

```
sacctmgr show user $USER format=user%20s,defaultaccount%30s
sacctmgr list association where
users=$USER format=account%30s,user%20s,qos%120s
```

Checking accounting share info - usage, fair-share, etc.

```
sshare -U
sshare -A $accountname
sshare -A $(sacctmgr -n show user $USER format=defaultaccount%30s)
sshare -a
```

Job launchers - basic (I)

```
#!/bin/bash -l
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH --time=0-00:05:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

echo "Hello from the batch queue on node ${SLURM_NODELIST}"
# Your more useful application can be started below!
```

Submit it with: sbatch launcher.sh

Job launchers - basic (II)

```
#!/bin/bash -l
#SBATCH -N 2
#SBATCH --ntasks-per-node=2
#SBATCH --time=0-03:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

echo "== Starting run at $(date)"
echo "== Job ID: ${SLURM_JOBID}"
echo "== Node list: ${SLURM_NODELIST}"
echo "== Submit dir. : ${SLURM_SUBMIT_DIR}"
# Your more useful application can be started below!
```

Job launchers - basic (III)

```
#!/bin/bash -l
#SBATCH -J MyTestJob
#SBATCH --mail-type=end,fail
#SBATCH --mail-user=Your.Email@Address.lu
#SBATCH -N 2
#SBATCH --ntasks-per-node=2
#SBATCH --time=0-03:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

echo "== Starting run at $(date)"
echo "== Job ID: ${SLURM_JOBID}"
echo "== Node list: ${SLURM_NODELIST}"
echo "== Submit dir. : ${SLURM_SUBMIT_DIR}"
# Your more useful application can be started below!
```

Job launchers - requesting memory

```
#!/bin/bash -l
#SBATCH -J MyLargeMemorySequentialJob
#SBATCH --mail-type=end,fail
#SBATCH --mail-user=Your.Email@Address.lu
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH --mem=64GB
#SBATCH --time=1-00:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

echo "== Starting run at $(date)"
echo "== Job ID: ${SLURM_JOBID}"
echo "== Node list: ${SLURM_NODELIST}"
echo "== Submit dir. : ${SLURM_SUBMIT_DIR}"
# Your more useful application can be started below!
```

Use "mem" to request (more) memory per node for low #core jobs

Job launchers - long jobs

```
#!/bin/bash -l
#SBATCH -J MyLongJob
#SBATCH --mail-type=all
#SBATCH --mail-user=Your.Email@Address.lu
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH --time=3-00:00:00
#SBATCH -p long
#SBATCH --qos=qos-long

echo "== Starting run at $(date)"
echo "== Job ID: ${SLURM_JOBID}"
echo "== Node list: ${SLURM_NODELIST}"
echo "== Submit dir. : ${SLURM_SUBMIT_DIR}"
# Your more useful application can be started below!
```

Longer walltime now possible but you should not (!) rely on it.
Always prefer batch and requeue-able jobs.

Job launchers - besteffort

```
#!/bin/bash -l
#SBATCH -J MyRerunnableJob
#SBATCH --mail-type=end,fail
#SBATCH --mail-user=Your.Email@Address.lu
#SBATCH -N 1
#SBATCH --ntasks-per-node=28
#SBATCH --time=0-12:00:00
#SBATCH -p batch
#SBATCH --qos=qos-besteffort
#SBATCH --requeue

echo "== Starting run at $(date)"
echo "== Job ID: ${SLURM_JOBID}"
echo "== Node list: ${SLURM_NODELIST}"
echo "== Submit dir. : ${SLURM_SUBMIT_DIR}"
# Your more useful application can be started below!
```

Many scientific applications support internal state saving and restart!
We will also discuss system-level checkpoint-restart with DMTCP.

Job launchers - threaded parallel

```
#!/bin/bash -l
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -c 28
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}
/path/to/your/threaded.app
```

By threaded we mean pthreads/OpenMP shared-memory applications.

Job launchers - MATLAB

```
#!/bin/bash -l
#SBATCH -N 1
#SBATCH --ntasks-per-node=28
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch
```

```
module load base/MATLAB
```

```
matlab -nodisplay -nosplash < /path/to/infile > /path/to/outfile
```

**MATLAB spawns processes, limited for now to single node execution.
We are still waiting for Distributed Computing Server availability.**

Job launchers - MATLAB

```
#!/bin/bash -l
#SBATCH -N 1
#SBATCH --ntasks-per-node=28
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch
```

```
module load base/MATLAB
```

```
matlab -nodisplay -nosplash < /path/to/infile > /path/to/outfile
```

**MATLAB spawns processes, limited for now to single node execution.
We are still waiting for Distributed Computing Server availability.**

A note on parallel jobs

**Currently the iris cluster is homogeneous.
Its core networking is a non-blocking fat-tree.**

- For now simply requesting a number of tasks (with 1 core/task) **should be performant**
- Different MPI implementations will however **behave differently**
 - ↪ very recent/latest versions available on **iris** for IntelMPI, OpenMPI, MVAPICH2
 - ↪ we ask that you let us know any perceived benefit for your applications when using one or the other
- We can make available optimized MPI-layer parameters obtained during our tuning runs
 - ↪ and hope they will improve even more your time to solution

Job launchers - IntelMPI

```
#!/bin/bash -l
#SBATCH -n 128
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

module load toolchain/intel
srun -n $SLURM_NTASKS /path/to/your/intel-toolchain-compiled-app
```

**IntelMPI is configured to use PMI2 for process management (optimal).
Bare mpirun will not work for now.**

Job launchers - OpenMPI

```
#!/bin/bash -l
#SBATCH -n 128
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

module load toolchain/foss
srun -n $SLURM_NTASKS /path/to/your/foss-toolchain-compiled-app
```

OpenMPI also uses PMI2 (again, optimal).
Bare mpirun does work but is not recommended.

You can easily generate a hostfile from within a SLURM job with:

```
srun hostname | sort -n > hostfile
```


Job launchers - MPI+OpenMP

```
#!/bin/bash -l
#SBATCH -N 10
#SBATCH --ntasks-per-node=1
#SBATCH -c 28
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

module load toolchain/intel
export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}
srun -n $SLURM_NTASKS /path/to/your/parallel-hybrid-app
```

Compile and use your applications in hybrid MPI+OpenMP mode when you can for better (best?) possible performance.



Summary

- 1 Introduction
- 2 SLURM workload manager
SLURM concepts and design for iris
Running jobs with SLURM
- 3 OAR and SLURM**
- 4 Conclusion

Notes on OAR

- OAR will remain the workload manager of Gaia and Chaos
 - ↪ celebrating **4250995** jobs on Gaia! (2017-11-07)
 - ↪ celebrating **1615659** jobs on Chaos! (2017-11-07)
- Many of its features are common to other workload managers, incl. SLURM
 - ↪ some things are exactly the same
 - ↪ but some things work in a different way
 - ↪ ... and some have no equivalent or are widely different
- An adjustment period for you is needed if you've only used OAR
 - ↪ next slides show [a brief transition guide](#)

OAR/SLURM - commands guide

Command	OAR (gaia/chaos)	SLURM (iris)
Submit passive/batch job	<code>oarsub -S \$script</code>	<code>sbatch \$script</code>
Start interactive job	<code>oarsub -I</code>	<code>srun -p interactive --qos qos-interactive --pty bash -i</code>
Queue status	<code>oarstat</code>	<code>squeue</code>
User job status	<code>oarstat -u \$user</code>	<code>squeue -u \$user</code>
Specific job status (detailed)	<code>oarstat -f -j \$jobid</code>	<code>scontrol show job \$jobid</code>
Delete (running/waiting) job	<code>oardele \$jobid</code>	<code>scancel \$jobid</code>
Hold job	<code>oarhold \$jobid</code>	<code>scontrol hold \$jobid</code>
Resume held job	<code>oarresume \$jobid</code>	<code>scontrol release \$jobid</code>
Node list and properties	<code>oarnodes</code>	<code>scontrol show nodes</code>

Similar yet different?

Many specifics will actually come from the way Iris is set up.

OAR/SLURM - job specifications

Specification	OAR	SLURM
Script directive	#OAR	#SBATCH
Nodes request	-l nodes=\$count	-N \$min-\$max
Cores request	-l core=\$count	-n \$count
Cores-per-node request	-l nodes=\$ncount/core=\$ccount	-N \$ncount --ntasks-per-node=\$ccount
Walltime request	-l [...],walltime=hh:mm:ss	-t \$min OR -t \$days-hh:mm:ss
Job array	--array \$count	--array \$specification
Job name	-n \$name	-J \$name
Job dependency	-a \$jobid	-d \$specification
Property request	-p "\$property=\$value"	-C \$specification

**Job specifications will need most adjustment on your side
... but thankfully Iris has a homogeneous configuration.
Running things in an optimal way will be much easier.**

OAR/SLURM - env. vars.

Environment variable	OAR	SLURM
Job ID	<code>\$OAR_JOB_ID</code>	<code>\$SLURM_JOB_ID</code>
Resource list	<code>\$OAR_NODEFILE</code>	<code>\$SLURM_NODELIST</code> #List not file! See below.
Job name	<code>\$OAR_JOB_NAME</code>	<code>\$SLURM_JOB_NAME</code>
Submitting user name	<code>\$OAR_USER</code>	<code>\$SLURM_JOB_USER</code>
Task ID within job array	<code>\$OAR_ARRAY_INDEX</code>	<code>\$SLURM_ARRAY_TASK_ID</code>
Working directory at submission	<code>\$OAR_WORKING_DIRECTORY</code>	<code>\$SLURM_SUBMIT_DIR</code>

Check available variables: `env | egrep "OAR|SLURM"`
Generate hostfile: `srun hostname | sort -n > hostfile`



Summary

- 1 Introduction
- 2 SLURM workload manager
SLURM concepts and design for iris
Running jobs with SLURM
- 3 OAR and SLURM
- 4 Conclusion

Conclusion and Practical Session start

We've discussed

- The design of SLURM for the **iris** cluster
- The permissions system in use through group accounts and QOS
- Main SLURM tools and how to use them
- Job types possible with SLURM on **iris**
- SLURM job launchers for sequential and parallel applications
- Transitioning from OAR to SLURM

And now..

Short DEMO time!

Conclusion and Practical Session start

We've discussed

- The design of SLURM for the **iris** cluster
- The permissions system in use through group accounts and QOS
- Main SLURM tools and how to use them
- Job types possible with SLURM on **iris**
- SLURM job launchers for sequential and parallel applications
- Transitioning from OAR to SLURM

And now..

Short DEMO time!

Your Turn!

Questions?

<http://hpc.uni.lu>

High Performance Computing @ UL

Prof. Pascal Bouvry

Dr. Sebastien Varrette & the UL HPC Team

(V. Plugaru, S. Peter, H. Cartiaux & C. Parisot)

University of Luxembourg, Belval Campus

Maison du Nombre, 4th floor

2, avenue de l'Université

L-4365 Esch-sur-Alzette

mail: hpc@uni.lu



1 Introduction

2 SLURM workload manager

SLURM concepts and design for iris
Running jobs with SLURM

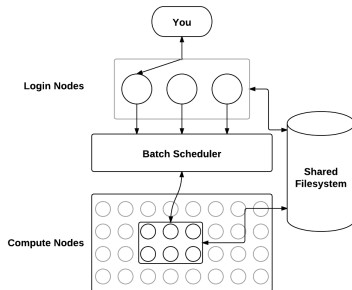
3 OAR and SLURM

4 Conclusion

Resource and Job Management Systems

• Resource and Job Management System (RJMS)

- “Glue” for a parallel computer to execute parallel jobs
- **Goal:** satisfy users' demands for computation
 - ✓ assign resources to user jobs with an efficient manner



Resource and Job Management Systems

• Resource and Job Management System (RJMS)

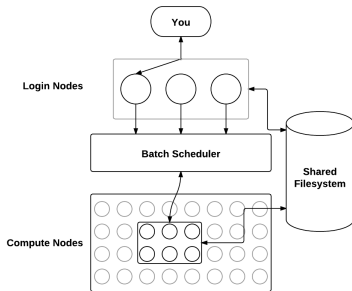
- “Glue” for a parallel computer to execute parallel jobs
- **Goal:** satisfy users’ demands for computation
 - ✓ assign resources to user jobs with an efficient manner

• HPC Resources:

- Nodes (typically a unique IP address)
 - ✓ Sockets / Cores / Hyperthreads
 - ✓ Memory
 - ✓ Interconnect/switch resources
- Generic resources (e.g. GPUs)
- Licenses

• Strategic Position

- Direct/constant knowledge of resources
- Launch and otherwise manage jobs



RJMS Layers

- Resource Allocation involves three principal abstraction layers:

- ↳ **Job Management:**

- ✓ declaration of a job & demand of resources and job characteristics,

- ↳ **Scheduling:** matching of the jobs upon the resources,

- ↳ **Resource Management:**

- ✓ launching and placement of job instances. . .
- ✓ . . . along with the job's control of execution

- **When there is more work than resources**

- ↳ the job scheduler manages queue(s) of work

- ✓ supports complex scheduling algorithms

- ↳ Supports resource limits (by queue, user, group, etc.)

RJMS Detailed Components

- **Resource Management**

- ↪ Resource Treatment (hierarchy, partitions,..)
- ↪ Job Launching, Propagation, Execution control
- ↪ Task Placement (topology, binding,..)
- ↪ **Advanced Features:**
 - ✓ High Availability, Energy Efficiency, Topology aware placement

RJMS Detailed Components

● Resource Management

- ↪ Resource Treatment (hierarchy, partitions,..)
- ↪ Job Launching, Propagation, Execution control
- ↪ Task Placement (topology, binding,..)
- ↪ **Advanced Features:**
 - ✓ High Availability, Energy Efficiency, Topology aware placement

● Job Management

- ↪ Job declaration and control (signaling, reprioritizing,...)
- ↪ Monitoring (reporting, visualization,..)
- ↪ **Advanced Features:**
 - ✓ Authentication (limitations, security,..)
 - ✓ QOS (checkpoint, suspend, accounting,...)
 - ✓ Interfacing (MPI libraries, debuggers, APIs,..)

RJMS Detailed Components

● Resource Management

- ↪ Resource Treatment (hierarchy, partitions,..)
- ↪ Job Launching, Propagation, Execution control
- ↪ Task Placement (topology, binding,..)
- ↪ **Advanced Features:**
 - ✓ High Availability, Energy Efficiency, Topology aware placement


● Job Management

- ↪ Job declaration and control (signaling, reprioritizing,...)
- ↪ Monitoring (reporting, visualization,..)
- ↪ **Advanced Features:**
 - ✓ Authentication (limitations, security,..)
 - ✓ QOS (checkpoint, suspend, accounting,...)
 - ✓ Interfacing (MPI libraries, debuggers, APIs,..)


● Scheduling

- ↪ Queues Management (priorities, multiple,..)
- ↪ Advanced Reservation


Job Scheduling



Job Scheduling (backfilling)



Job Scheduling (suspension & requeue)



Main Job Schedulers

Name	Company	Version*
SLURM	SchedMD	17.02.8
LSF	IBM	10.1
OpenLava	LSF Fork	2.2
MOAB/Torque	Adaptative Computing	6.1
PBS	Altair	13.0
OAR (PBS Fork)	LIG	2.5.7
Oracle Grid Engine (formerly SGE)	Oracle	

*: As of Oct. 2017

Main Job Schedulers

Name	Company	Version*
SLURM	SchedMD	17.02.8
LSF	IBM	10.1
OpenLava	LSF Fork	2.2
MOAB/Torque	Adaptative Computing	6.1
PBS	Altair	13.0
OAR (PBS Fork)	LIG	2.5.7
Oracle Grid Engine (formely SGE)	Oracle	

*: As of Oct. 2017

UL HPC resource manager: OAR

The OAR Batch Scheduler

<http://oar.imag.fr>

- Versatile resource and task manager

- ↳ schedule **jobs** for users on the cluster **resource**
- ↳ OAR resource = a node or part of it (CPU/core)
- ↳ OAR job = execution time (**walltime**) on a set of resources



UL HPC resource manager: OAR

The OAR Batch Scheduler

<http://oar.imag.fr>

- Versatile resource and task manager
 - ↪ schedule **jobs** for users on the cluster **resource**
 - ↪ OAR resource = a node or part of it (CPU/core)
 - ↪ OAR job = execution time (**walltime**) on a set of resources



OAR main features includes:

- **interactive vs. passive (aka. batch) jobs**
- **best effort jobs**: use more resource, accept their release any time
- **deploy jobs (Grid5000 only)**: deploy a customized OS environment
 - ↪ ... and have full (root) access to the resources
- **powerful resource filtering/matching**

Main OAR commands

oarsub submit/reserve a job (by default: **1 core for 2 hours**)

oardel delete a submitted job

oarnodes shows the resources states

oarstat shows information about running or planned jobs

	Submission
interactive	<code>oarsub [options] -I</code>
passive	<code>oarsub [options] scriptName</code>

- Each created job receive an identifier JobID
↪ Default passive job log files: `OAR.JobID.std{out,err}`
- You can make a reservation with `-r "YYYY-MM-DD HH:MM:SS"`

Main OAR commands

- oarsub** submit/reserve a job (by default: **1 core for 2 hours**)
- oardel** delete a submitted job
- oarnodes** shows the resources states
- oarstat** shows information about running or planned jobs

	Submission
interactive	<code>oarsub [options] -I</code>
passive	<code>oarsub [options] scriptName</code>

- Each created job receive an identifier JobID
 - ↪ Default passive job log files: `OAR.JobID.std{out,err}`
- You can make a reservation with `-r "YYYY-MM-DD HH:MM:SS"`

Direct access to nodes by `ssh` is forbidden: use `oarsh` instead

OAR job environment variables

Once a job is created, some environments variables are defined:

Variable	Description
\$OAR_NODEFILE	Filename which lists all reserved nodes for this job
\$OAR_JOB_ID	OAR job identifier
\$OAR_RESOURCE_PROPERTIES_FILE	Filename which lists all resources and their properties
\$OAR_JOB_NAME	Name of the job given by the "-n" option of oarsub
\$OAR_PROJECT_NAME	Job project name

Useful for MPI jobs for instance:

```
$> mpirun -machinefile $OAR_NODEFILE /path/to/myprog
```

... Or to collect how many cores are reserved per node:

```
$> cat $OAR_NODEFILE | uniq -c
```

OAR job types

Job Type	Max Walltime (hour)	Max #active_jobs	Max #active_jobs_per_user
interactive	12:00:00	10000	5
default	120:00:00	30000	10
besteffort	9000:00:00	10000	1000

cf /etc/oar/admission_rules/*.conf

- **interactive**: useful to test / prepare an experiment
 - ↪ you get a shell on the first reserved resource
- **best-effort vs. default**: nearly unlimited constraints **YET**
 - ↪ a besteffort job can be killed as soon as a default job as no other place to go
 - ↪ enforce checkpointing (and/or idempotent) strategy

Characterizing OAR resources

Specifying wanted resources in a hierarchical manner

- Use the `-l` option of `oarsub`. Main constraints:

<code>enclosure=N</code>	number of enclosure
<code>nodes=N</code>	number of nodes
<code>core=N</code>	number of cores
<code>walltime=hh:mm:ss</code>	job's max duration

Specifying OAR resource properties

- Use the `-p` option of `oarsub`:

Syntax: `-p "property='value'"`

<code>gpu='{YES,NO}'</code>	has (or not) a GPU card
<code>host='fqdn'</code>	full hostname of the resource
<code>network_address='hostname'</code>	Short hostname of the resource
(Chaos only) <code>nodeclass='{k,b,h,d,r}'</code>	Class of node

OAR (interactive) job examples

- 2 cores on 3 nodes (same enclosure) for 3h15:

Total: 6 cores

```
(frontend)$> oarsub -I -l /enclosure=1/nodes=3/core=2,walltime=3:15
```

OAR (interactive) job examples

- 2 cores on 3 nodes (same enclosure) for 3h15:

Total: 6 cores

```
(frontend)$> oarsub -I -l /enclosure=1/nodes=3/core=2,walltime=3:15
```

- 4 cores on a GPU node for 8 hours

Total: 4 cores

```
(frontend)$> oarsub -I -l /core=4,walltime=8 -p "gpu='YES'"
```

OAR (interactive) job examples

- 2 cores on 3 nodes (same enclosure) for 3h15: Total: 6 cores

```
(frontend)$> oarsub -I -l /enclosure=1/nodes=3/core=2,walltime=3:15
```

- 4 cores on a GPU node for 8 hours Total: 4 cores

```
(frontend)$> oarsub -I -l /core=4,walltime=8 -p "gpu='YES'"
```

- 2 nodes among the h-cluster1-* nodes (Chaos only) Total: 24 cores

```
(frontend)$> oarsub -I -l nodes=2 -p "nodeclass='h'"
```

OAR (interactive) job examples

- 2 cores on 3 nodes (same enclosure) for 3h15: Total: 6 cores

```
(frontend)$> oarsub -I -l /enclosure=1/nodes=3/core=2,walltime=3:15
```

- 4 cores on a GPU node for 8 hours Total: 4 cores

```
(frontend)$> oarsub -I -l /core=4,walltime=8 -p "gpu='YES'"
```

- 2 nodes among the h-cluster1-* nodes (Chaos only) Total: 24 cores

```
(frontend)$> oarsub -I -l nodes=2 -p "nodeclass='h'"
```

- 4 cores on 2 GPU nodes + 20 cores on other nodes Total: 28 cores

```
$> oarsub -I -l "{gpu='YES'}/nodes=2/core=4+{gpu='NO'}/core=20"
```


OAR (interactive) job examples

- 2 cores on 3 nodes (same enclosure) for 3h15: Total: 6 cores

```
(frontend)$> oarsub -I -l /enclosure=1/nodes=3/core=2,walltime=3:15
```

- 4 cores on a GPU node for 8 hours Total: 4 cores

```
(frontend)$> oarsub -I -l /core=4,walltime=8 -p "gpu='YES'"
```

- 2 nodes among the h-cluster1-* nodes (Chaos only) Total: 24 cores

```
(frontend)$> oarsub -I -l nodes=2 -p "nodeclass='h'"
```

- 4 cores on 2 GPU nodes + 20 cores on other nodes Total: 28 cores

```
$> oarsub -I -l "{gpu='YES'}/nodes=2/core=4+{gpu='NO'}/core=20"
```

- A full big SMP node Total: 160 cores on gaia-74

```
$> oarsub -t bigsmp -I 1 node=1
```

Some other useful features of OAR

Connect to a running job

```
(frontend)$> oarsub -C JobID
```

Cancel a job

```
(frontend)$> oardel JobID
```

Status of a jobs

```
(frontend)$> oarstat -state -j JobID
```

View the job

```
(frontend)$> oarstat
```

```
(frontend)$> oarstat -f -j JobID
```

Get info on the nodes

```
(frontend)$> oarnodes
```

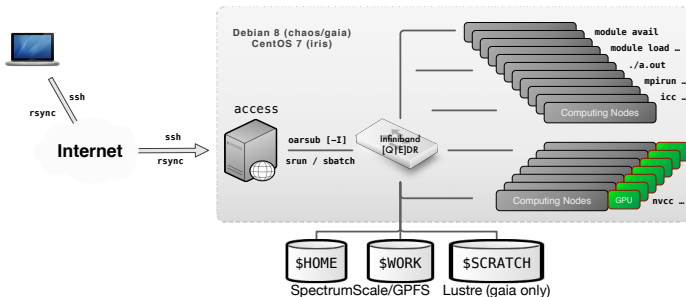
```
(frontend)$> oarnodes -l
```

```
(frontend)$> oarnodes -s
```

Run a best-effort job

```
(frontend)$> oarsub -t besteffort ...
```

OAR Practical session



Demo Time

- gaia or chaos UL cluster access
- Interactive / Passive job submission

Designing efficient OAR job launchers

Resources/Example

<https://github.com/ULHPC/launcher-scripts>



- UL HPC grant access to **parallel computing** resources
 - ideally: OpenMP/MPI/CUDA/OpenCL jobs
 - if serial jobs/tasks: run them efficiently
- Avoid to submit purely serial jobs to the OAR queue a
 - waste the computational power (11 out of 12 cores on gaia).
 - use whole nodes by running at least 12 serial runs at once
- **Key:** understand difference between **Task** and **OAR job**

Designing efficient OAR job launchers

Resources/Example

<https://github.com/ULHPC/launcher-scripts>



- UL HPC grant access to **parallel computing** resources
 - ↪ ideally: OpenMP/MPI/CUDA/OpenCL jobs
 - ↪ if serial jobs/tasks: run them efficiently
- Avoid to submit purely serial jobs to the OAR queue a
 - ↪ waste the computational power (11 out of 12 cores on gaia).
 - ↪ use whole nodes by running at least 12 serial runs at once
- **Key:** understand difference between **Task** and **OAR job**

For more information...

- Incoming Practical Session
 - ↪ HPC workflow with sequential jobs (C,python,java etc.)

OAR Simple Example of usage

```
# Simple interactive job
(access)$> oarsub -I
(node)$> echo $OAR_JOBID
4239985
(node)$> echo $OAR_NODEFILE
/var/lib/oar//4239985
(node)$> cat $OAR_NODEFILE | wc -l
8
(node)$> cat $OAR_NODEFILE
moonshot1-39
moonshot1-39
moonshot1-39
moonshot1-39
moonshot1-40
moonshot1-40
moonshot1-40
moonshot1-40
```

View existing job

```
# View YOUR jobs (remove -u to view all)
```

```
(access)$> oarstat -u
```

Job id	Name	User	Submission Date	S	Queue
4239985		svarrette	2017-10-23 12:33:41	R	default

View Detailed info on jobs

```
(access)$> oarstat -f -j 4239985
```

```
Job_Id: 4239985
```

```
[...]
```

```
state = Running
```

```
wanted_resources = -l "{type = 'default'}/ibpool=1/host=2,walltime=2:0:0
```

```
types = interactive, inner=4236343, moonshot
```

```
assigned_resources = 3309+3310+3311+3312+3313+3314+3315+3316
```

```
assigned_hostnames = moonshot1-39+moonshot1-40
```

```
queue = default
```

```
launchingDirectory = /home/users/svarrette
```

```
stdout_file = OAR.4239985.stdout
```

```
stderr_file = OAR.4239985.stderr
```

```
jobType = INTERACTIVE
```

```
properties = (((bigmem='NO' AND bigsmp='NO') AND dedicated='NO') AND os=
```

```
walltime = 2:0:0
```

```
initial_request = oarsub -I -l nodes=2 -t moonshot -t inner=4236343
```

```
message = R=8,W=2:0:0,J=I,T=inner|interactive|moonshot (Karma=1.341)
```


Access to an existing job: Attempt 1

```
# Get your job ID...  
(access)$> oarstat -u
```

- **Attempt 1:** Get assigned resources and ...

Access to an existing job: Attempt 1

```
# Get your job ID...  
(access)$> oarstat -u
```

- **Attempt 1:** Get assigned resources and ... ssh to it!

```
# Collect the assigned resources  
(access)$> oarstat -f -j 4239985 | grep hostname  
    assigned_hostnames = moonshot1-39+moonshot1-40
```

```
(access)$> ssh moonshot1-39  
[...]
```

```
=====  
/!\ WARNING:   Direct login by ssh is forbidden.
```

Use `oarsub(1)` to reserve nodes, and `oarsh(1)` to connect to your reserved nodes, typically by:

```
    OAR_JOB_ID=<jobid> oarsh <nodename>
```

Access to an existing job

- Using oarsh:

```
# Get your job ID...  
(access)$> oarstat -u  
# ... get the hostname of the nodes allocated ...  
(access)$> oarstat -f -j 4239985 | grep hostname  
# ... and connect to it with oarsh  
(access)$> OAR_JOB_ID=4239985 oarsh moonshot1-39
```

Access to an existing job

- Using oarsh:

```
# Get your job ID...  
(access)$> oarstat -u  
# ... get the hostname of the nodes allocated ...  
(access)$> oarstat -f -j 4239985 | grep hostname  
# ... and connect to it with oarsh  
(access)$> OAR_JOB_ID=4239985 oarsh moonshot1-39
```

- (better) Using oarsub -C:

```
# Get your job ID...  
(access)$> oarstat -u  
# ... and connect to the FIRST node of the reservation  
(access)$> oarsub -C 4239985
```

MPI jobs

- Intel MPI

```
(node)$> module load toolchain/intel  
# ONLY on moonshot node have no IB card: export I_MPI_FABRICS=tcp  
(node)$> mpirun -hostfile $OAR_NODEFILE /path/to/mpiprogram
```

- OpenMPI:

```
(node)$> module load mpi/OpenMPI  
(node)$> mpirun -hostfile $OAR_NODEFILE -x PATH -x LD_LIBRARY_PATH \  
/path/to/mpiprogram
```

- For more details: See MPI sessions

https://github.com/ULHPC/launcher-scripts/blob/devel/bash/MPI/mpi_launcher.sh

OAR Launcher Scripts

```
$> oarsub -S <scriptname>           # -S: interpret #OAR comments
```

- Our launcher scripts on Github: <https://github.com/ULHPC/launcher-scripts>
↳ see in particular [our generic launcher](#) compliant w. OAR & SLURM
- **Example:**

```
#!/bin/bash
#OAR -l nodes=2/core=1,walltime=1
#OAR -n MyNamedJob
# Prepare UL HPC modules
if [ -f /etc/profile ]; then
. /etc/profile
fi

module load toolchain/intel
/path/to/prog <ARGS>
```

Slurm Workload Manager

- **Simple Linux Utility for Resource Management**
 - ↪ Development started in 2002
 - ✓ initially as a simple resource manager for Linux clusters
 - ↪ Has evolved into a capable job scheduler through use of opt. plugins
 - ↪ About 500,000 lines of C code today.
 - ↪ Supports AIX, Linux, Solaris, other Unix variants
- Used on many of the world's largest computers

Slurm Workload Manager

- **Simple Linux Utility for Resource Management**
 - ↪ Development started in 2002
 - ✓ initially as a simple resource manager for Linux clusters
 - ↪ Has evolved into a capable job scheduler through use of opt. plugins
 - ↪ About 500,000 lines of C code today.
 - ↪ Supports AIX, Linux, Solaris, other Unix variants
 - Used on many of the world's largest computers
-
- **Now deployed on new UL HPC clusters**
 - ↪ starting iris cluster (2017)

Slurm Design Goals

- **Small and simple**
 - **Highly scalable** and **Fast**
 - ↪ managing 1.6 million core IBM BlueGene/Q,
 - ↪ tested to 33 million cores using emulation
 - ↪ throughput: up to 600 jobs p.s. & 1000 job submissions p.s.
 - **Modular:**
 - ↪ plugins to support \neq scheduling policies, MPI librairies. . .
 - **Secure** and **Fault-tolerant**
 - ↪ highly tolerant of system failures
 - **Power Management** and detailed monitoring
-
- **Open source:** GPL v2, active world-wide development
 - **Portable:** written in C with a GNU autoconf configuration engine

Slurm Docs and Resources

- User and Admins latest documentation:
↳ <http://slurm.schedmd.com/documentation.html>
- Detailed man pages for commands and configuration files
↳ http://slurm.schedmd.com/man_index.html
- All SLURM related publications and presentations:
↳ <http://slurm.schedmd.com/publications.html>

ULHPC Documentation & comparison to OAR

<https://hpc.uni.lu/users/docs/scheduler.html>

Other Resources

• Puppet module ULHPC/slurm

- Developed by the UL HPC Team – see <https://forge.puppet.com/ULHPC/>
- Used in production on iris cluster
- see also Slurm Control Repo Example

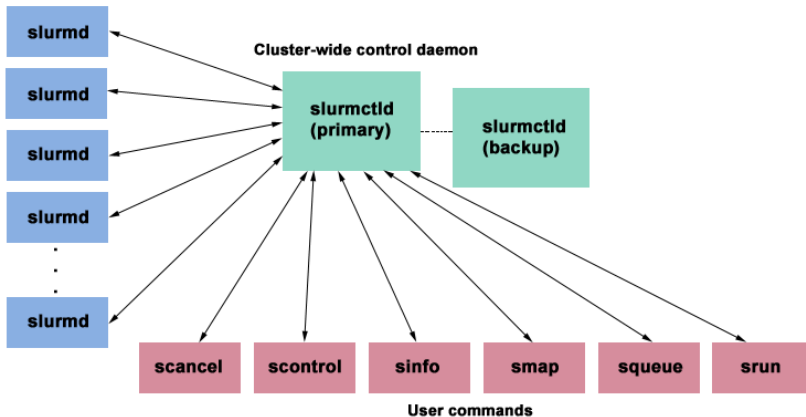


Puppet Class	Description
<code>slurm</code>	The main slurm class, piloting all aspects of the configuration
<code>slurm::slurmdbd</code>	Specialized class for <code>Slurmdbd</code> , the Slurm Database Daemon.
<code>slurm::slurmctld</code>	Specialized class for <code>Slurmctld</code> , the central management daemon of Slurm.
<code>slurm::slurmd</code>	Specialized class for <code>Slurmd</code> , the compute node daemon for Slurm.
<code>slurm::munge</code>	Manages <code>MUNGE</code> , an authentication service for creating and validating credentials.
<code>slurm::pam</code>	Handle PAM aspects for SLURM (Memlock for MPI etc.)

Puppet Defines	Description
<code>slurm::download</code>	takes care of downloading the SLURM sources for a given version passed as resource name
<code>slurm::build</code>	building Slurm sources into packages (_i.e. RPMs for the moment)
<code>slurm::install::packages</code>	installs the Slurm packages, typically built from <code>slurm::build</code>
<code>slurm::acct::*</code>	adding (or removing) accounting resources to the slurm database


SLURM Architecture

One daemon per node

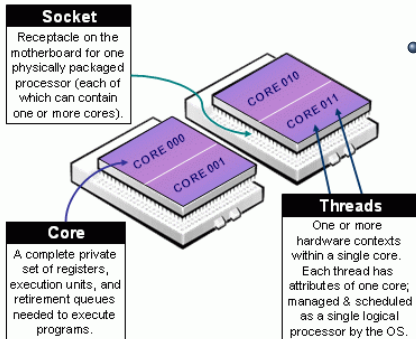


SLURM Entities

Entity	Description
Computing node	Computer used for the execution of programs
Partition	Group of nodes into logical sets
Job	Allocation of resources assigned to a user for some time
Job Step	Sets of (possible parallel) tasks with a job




SLURM Multi-Core/Thread Support



Nodes hierarchy


- ↪ NUMA [base]board
 - ✓ Socket/Core/Thread
- ↪ Memory
- ↪ Generic Resources GRES (e.g. GPUs)

SLURM Entities example



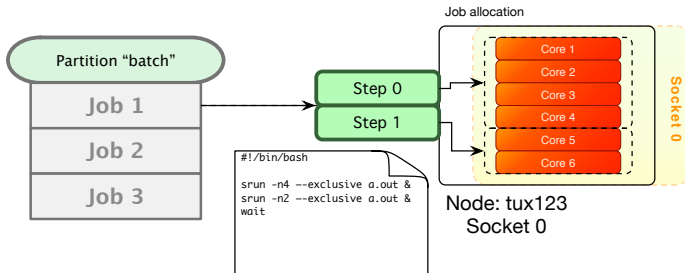
- Users submit jobs to a partition (queue)

SLURM Entities example



- Jobs are allocated resources

SLURM Entities example



- Jobs spawn steps, which are allocated resources from within the job's allocation

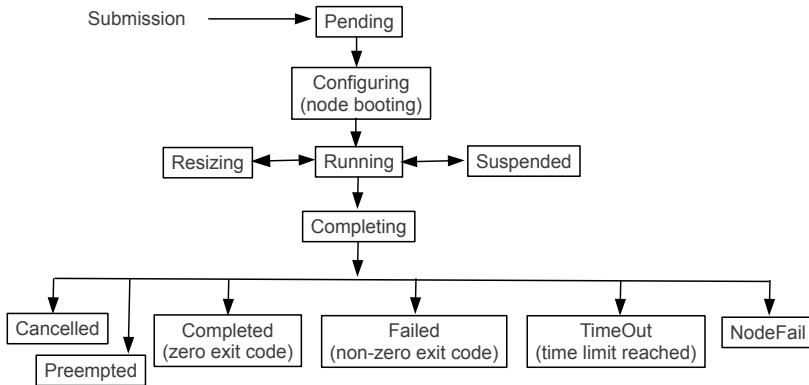
Node State Information

- NUMA boards, Sockets, Cores, Threads
- CPUs
 - ↪ can treat each core or each thread as a CPU for scheduling purposes
- Memory size
- Temporary disk space
- Features (arbitrary string, e.g. OS version)
- Weight (scheduling priority, . . .)
 - ↪ can favor least capable node that satisfies job requirement
- Boot time
- CPU Load
- State (e.g. drain, down, etc.)
- Reason, time and user ID
 - ↪ e.g. "Bad PDU [operator@12:40:10T12/20/2011]"

Queues/Partitions State Information

- Associated with specific set of nodes
 - ↳ Nodes can be in more than one partition (not the case in `iris`)
- Job size and time limits
- Access control list (by Linux group) / QoS
- Preemption rules
- State information (e.g. drain)
- Over-subscription and gang scheduling rules

Job State



Slurm Daemons

- **slurmctld**: Central controller (typically one per cluster)
 - ↪ Optional backup with automatic fail over
 - ↪ Monitors state of resources
 - ↪ Manages job queues and Allocates resources
- **slurmd**: Compute node daemon
 - ↪ typically one per compute node, one or more on front-end nodes
 - ↪ Launches and manages tasks
 - ↪ Small and very light-weight (low memory and CPU use)

- **Common configuration file**: `/etc/slurm/slurm.conf`
 - ↪ Other interesting files: `/etc/slurm/{topology,gres}.conf`

- **slurmdbd**: database daemon (typically one per site)
 - ↪ Collects accounting information
 - ↪ Uploads configuration information (limits, fair-share, etc.)

Slurm in iris cluster

- Predefined **Queues/Partitions**:
 - ↪ batch (Default) **Max**: 30 nodes, 5 days walltime
 - ↪ interactive **Max**: 2 nodes, 4h walltime, 10 jobs
 - ↪ long **Max**: 2 nodes, 30 days walltime, 10 jobs
- Corresponding Quality of Service (QOS)
- Possibility to run besteffort jobs via the qos-besteffort QOS
- Accounts associated to supervisor (multiple associations possible)
- Proper group/user accounting

Slurm Job Management

- User jobs have the following key characteristics:
 - ↪ set of requested resources:
 - ✓ number of computing resources: **nodes** (including all their CPUs and cores) or **CPUs** (including all their cores) or **cores**
 - ✓ amount of **memory**: either per node or per CPU
 - ✓ **(wall)time** needed for the user's tasks to complete their work
 - ↪ a requested node **partition** (job queue)
 - ↪ a requested **quality of service** (QoS) level which grants users specific accesses
 - ↪ a requested **account** for accounting purposes

By default...

- users submit jobs to a particular partition, **and** under a particular account (pre-set per user).

Slurm Commands: General Info

- Man pages available for all commands, daemons and config. files
 - ↪ `--help` option prints brief description of all options
 - ↪ `--usage` option prints a list of the options
 - ↪ `-v` | `-vv` | `-vvv`: verbose output
- Commands can be run on any node in the cluster
- Any failure results in a non-zero exit code
- APIs make new tool development easy
 - ↪ Man pages available for all APIs
- Almost all options have two formats
 - ↪ A single letter option (e.g. `-p` batch for partition 'batch')
 - ↪ A verbose option (e.g. `--partition=batch`)
- Time formats: `DD-HH:MM::SS`

User Commands: Job/step Allocation

- **sbatch**: Submit script for later execution (batch mode)
 - ↪ allocate resources (nodes, tasks, partition, etc.)
 - ↪ Launch a script containing `sruns` for series of steps on them.
- **salloc**: Create job allocation & start a shell to use it
 - ↪ allocate resources (nodes, tasks, partition, etc.),
 - ↪ either run a command or start a shell.
 - ↪ Request launch `srun` from shell. (interactive commands within one allocation)
- **srunk**: Create a job allocation (if needed) and launch a job step (typically an MPI job)
 - ↪ allocate resources (number of nodes, tasks, partition, constraints, etc.)
 - ↪ launch a job that will execute on them.
- **sattach**: attach to running job for debuggers



User & Admin Commands: System Information

- **sinfo**: Report system status (nodes, partitions etc.)
- **squeue**: display jobs[steps] and their state
- **scancel**: cancel a job or set of jobs.
- **scontrol**: view and/or update system, nodes, job, step, partition or reservation status
- **sstat**: show status of running jobs.
- **sacct**: display accounting information on jobs.
- **sprio**: show factors that comprise a jobs scheduling priority
- **smap**: graphically show information on jobs, nodes, partitions
 - ↪ not available on iris

Slurm Admin Commands

- **sacctmgr**: setup accounts, specify limitations on users and groups.
- **sshare**: view sharing information from multifactor plugin.
- **sreport**: display information from accounting database on jobs, users, clusters.
- **sview**: graphical view of cluster. Display and change characteristics of jobs, nodes, partitions.
 - ↳ not yet available on iris cluster
- **strigger**: show, set, clear event triggers. Events are usually system events such as an equipment failure.

Slurm vs. OAR Main Commands

Action	SLURM command	OAR Command
Submit passive/batch job	<code>sbatch [...] \$script</code>	<code>oarsub [...] \$script</code>
Start interactive job	<code>srun [...] --pty bash</code>	<code>oarsub -I [...]</code>
Queue status	<code>squeue</code>	<code>oarstat</code>
User job status	<code>squeue -u \$user</code>	<code>oarstat -u \$user</code>
Specific job status (detailed)	<code>scontrol show job \$jobid</code>	<code>oarstat -f -j \$jobid</code>
Job accounting status (detailed)	<code>sacct --job \$jobid -l</code>	
Delete (running/waiting) job	<code>scancel \$jobid</code>	<code>oardel \$jobid</code>
Hold job	<code>scontrol hold \$jobid</code>	<code>oarhold \$jobid</code>
Resume held job	<code>scontrol release \$jobid</code>	<code>oarresume \$jobid</code>
Node list and their properties	<code>scontrol show nodes</code>	<code>oarnodes</code>

Job Specifications

Specification	SLURM	OAR
Script directive	#SBATCH	#OAR
<n> Nodes request	-N <n>	-l nodes=<n>
<n> Cores/Tasks request	-n <n>	-l core=<n>
<c> Cores-per-node request	--ntasks-per-node=<c>	-l nodes=<n>/core=<c>
<c> Cores-per-task request (multithreading)	-c=<c>	
<m> GB memory per node request	--mem=<m>GB	
Walltime request	-t <mm>/<days>-hh[:mm:ss]>	-l walltime=hh[:mm:ss]
Job array	--array <specification>	--array <count>
Job name	-J <name>	-n <name>
Job dependency	-d <specification>	-a <jobid>
Property request	-C <specification>	-p "<property>=<value>"
Specify job partition/queue	-p <partition>	-t <queue>
Specify job qos	--qos <qos>	
Specify account	-A <account>	
Specify email address	--mail-user=<email>	--notify "mail:<email>"

Typical Workflow

```
# Run an interactive job -- make an alias 'si [...]'
$> srun -p interactive --qos qos-interactive --pty bash
# Ex: interactive job for 30 minutes, with 2 nodes/4 tasks per node
$> si --time=0:30:0 -N 2 --ntasks-per-node=4
# Run a [passive] batch job -- make an alias 'sb [...]'
$> sbatch -p batch --qos qos-batch /path/to/launcher.sh
# Will create (by default) slurm-<jobid>.out file
```

Environment variable	SLURM	OAR
Job ID	<code>\$SLURM_JOB_ID</code>	<code>\$OAR_JOB_ID</code>
Resource list	<code>\$SLURM_NODELIST</code> #List not file!	<code>\$OAR_NODEFILE</code>
Job name	<code>\$SLURM_JOB_NAME</code>	<code>\$OAR_JOB_NAME</code>
Submitting user name	<code>\$SLURM_JOB_USER</code>	<code>\$OAR_USER</code>
Task ID within job array	<code>\$SLURM_ARRAY_TASK_ID</code>	<code>\$OAR_ARRAY_INDEX</code>
Working directory at submission	<code>\$SLURM_SUBMIT_DIR</code>	<code>\$OAR_WORKING_DIRECTORY</code>
Number of nodes assigned to the job	<code>\$SLURM_NNODES</code>	
Number of tasks of the job	<code>\$SLURM_NTASKS</code>	<code>\$(wc -l \${OAR_NODEFILE})</code>

- **Note:** create the equivalent of `$OAR_NODEFILE` in Slurm:

```
↪ srun hostname | sort -n > hostfile
```

Available Node partitions

- Slurm Command Option `-p, --partition=<partition>`
↳ **Ex:** `{srun,sbatch} -p batch [...]`
- Date format: `-t <minutes>` or `-t <D>-<H>:<M>:<S>`

Partition	#Nodes	Default time	Max time	Max nodes/user
batch	80%	0-2:0:0 [2h]	5-0:0:0 [5d]	unlimited
interactive	10%	0-1:0:0 [1h]	0-4:0:0 [4h]	2
long	10%	0-2:0:0 [2h]	30-0:0:0 [30d]	2

Quality of Service (QoS)

- Slurm Command Option `--qos=<qos>`
- **There is no default QOS** (due to the selected scheduling model)
 - you **MUST** provide upon **any job submission**
 - a default qos is guessed from the partition **i.e.** `qos-<partition>`

QoS	User group	Max cores	Max jobs/user	Description
qos-besteffort	ALL	no limit		Preemptible jobs, queued on preemption
qos-batch	ALL	1064	100	Normal usage of the batch partition
qos-interactive	ALL	224	10	Normal usage of the interactive partition
qos-long	ALL	224	10	Normal usage of the long partition
qos-batch-###	rsvd	rsvd	100	Reserved usage of the batch partition
qos-interactive-###	rsvd	rsvd	10	Reserved usage of the interactive partition
qos-long-###	rsvd	rsvd	10	Reserved usage of the long partition

Accounts

- Every user job runs under a group account
 ↳ granting access to specific QOS levels.

Account	Parent Account
UL	
FSTC	UL
FDEF	UL
FLSHASE	UL
LCSB	UL
SNT	UL
Professor \$X	<i>FACULTY/IC</i>
Group head \$G	<i>FACULTY/IC</i>
Researcher \$R	Professor \$X
Researcher \$R	Group head \$G
Student \$S	Professor \$X
Student \$S	Group head \$G
External collaborator \$E	Professor \$X
External collaborator \$E	Group head \$G

```
$> sacctmgr list associations where users=$USER \
      format=Account%30s,User,Partition,QOS
```

Other Features

- **Checkpoint / Restart**

- ↪ Based on DMTCP: Distributed MultiThreaded CheckPointing
- ↪ see the official DMTCP launchers
- ↪ ULHPC example

- Many metrics can be extracted from user jobs

- ↪ with SLURM's own tools (sacct/sstat)
- ↪ within the jobs with e.g. **PAPI**
- ↪ easy to bind executions with Allinea Performance Report

- Advanced admission rules

- ↪ to simplify CLI

- Container **Shifter** / **Singularity**

- ↪ Work in progress, not yet available on **iris**

Simple Example of usage

```
# Simple interactive job  
(access)$> srun -p interactive [--qos qos-interactive] --pty bash  
(node)$> echo $SLURM_NTASKS  
1  
(node)$> echo $SLURM_JOBID  
59900
```

Simple Example of usage

```
# Simple interactive job
(access)$> srun -p interactive [--qos qos-interactive] --pty bash
(node)$> echo $SLURM_NTASKS
1
(node)$> echo $SLURM_JOBID
59900
```

```
$> squeue -u $USER -l    # OR 'sq'
```

Simple Example of usage

```
# Simple interactive job
(access)$> srun -p interactive [--qos qos-interactive] --pty bash
(node)$> echo $SLURM_NTASKS
1
(node)$> echo $SLURM_JOBID
59900
```

```
$> squeue -u $USER -l    # OR 'sq'
```

- Many metrics during (scontrol)/after job execution (sacct)
 - ↳ including energy (J) – but with caveats
 - ↳ job **steps** counted individually
 - ↳ enabling advanced application debugging and optimization
- Job information available in easily parseable format (add -p/-P)

Live Job Statistics

```
$> scontrol show job 59900
JobId=59900 JobName=bash
UserId=<login>(<uid>) GroupId=clusterusers(666) MCS_label=N/A
Priority=6627 Nice=0 Account=ulhpc QOS=qos-interactive
JobState=RUNNING Reason=None Dependency=(null)
RunTime=00:04:19 TimeLimit=01:00:00 TimeMin=N/A
SubmitTime=2017-10-22T23:07:02 EligibleTime=2017-10-22T23:07:02
StartTime=2017-10-22T23:07:02 EndTime=2017-10-23T00:07:02 Deadline=N/A
PreemptTime=None SuspendTime=None SecsPreSuspend=0
Partition=interactive AllocNode:Sid=access1:72734
[...]
NodeList=iris-002
NumNodes=1 NumCPUs=1 NumTasks=1 CPUs/Task=1 ReqB:S:C:T=0:0:*:*
TRES=cpu=1,mem=4G,node=1
Socks/Node=* NtasksPerN:B:S:C=0:0:*:* CoreSpec=*
MinCPUsNode=1 MinMemoryCPU=4G MinTmpDiskNode=0
[...]
Command=bash
WorkDir=/mnt/irisgpfs/users/<login>
```

Node/Job Statistics

```
$> sinfo
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
interactive	up	4:00:00	10	idle	iris-[001-010]
long	up	30-00:00:0	2	resv	iris-[019-020]
long	up	30-00:00:0	8	idle	iris-[011-018]
batch*	up	5-00:00:00	5	mix	iris-[055,060-062,101]
batch*	up	5-00:00:00	13	alloc	iris-[053-054,056-059,102-108]
batch*	up	5-00:00:00	70	idle	iris-[021-052,063-100]

Node/Job Statistics

```
$> sinfo
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
interactive	up	4:00:00	10	idle	iris-[001-010]
long	up	30-00:00:0	2	resv	iris-[019-020]
long	up	30-00:00:0	8	idle	iris-[011-018]
batch*	up	5-00:00:00	5	mix	iris-[055,060-062,101]
batch*	up	5-00:00:00	13	alloc	iris-[053-054,056-059,102-108]
batch*	up	5-00:00:00	70	idle	iris-[021-052,063-100]

```
$> sacct --format=account,user,jobid,jobname,partition,state -j <JOBID>
```

```
$> sacct --format=elapsed,elapseddraw,start,end -j <JOBID>
```

```
$> sacct -j <JOBID> --format=maxrss,maxvmsize,consumedenergy,\
      consumedenergyraw,nodelist
```


Playing with hostname and task ID label

```
$> srun [-N #node] [-n #task] [--ntasks-per-node #n] [] CMD
```

```
# -n: #tasks
```

```
$> srun -n 4 -l hostname
```

```
1: iris-055
```

```
2: iris-055
```

```
3: iris-055
```

```
0: iris-055
```

Playing with hostname and task ID label

```
$> srun [-N #node] [-n #task] [--ntasks-per-node #n] [] CMD
```

```
# -n: #tasks
```

```
$> srun -n 4 -l hostname
```

```
1: iris-055
```

```
2: iris-055
```

```
3: iris-055
```

```
0: iris-055
```

```
# -N: #nodes
```

```
$> srun -N 4 -l hostname
```

```
3: iris-058
```

```
2: iris-057
```

```
1: iris-056
```

```
0: iris-055
```

Playing with hostname and task ID label

```
$> srun [-N #node] [-n #task] [--ntasks-per-node #n] [] CMD
```

-n: #tasks

```
$> srun -n 4 -l hostname
```

```
1: iris-055
```

```
2: iris-055
```

```
3: iris-055
```

```
0: iris-055
```

-c: #cpus/task ~#thread/task

```
$> srun -c 4 -l hostname
```

```
0: iris-055
```

-N: #nodes

```
$> srun -N 4 -l hostname
```

```
3: iris-058
```

```
2: iris-057
```

```
1: iris-056
```

```
0: iris-055
```

Playing with hostname and task ID label

```
$> srun [-N #node] [-n #task] [--ntasks-per-node #n] [] CMD
```

-n: #tasks

```
$> srun -n 4 -l hostname  
1: iris-055  
2: iris-055  
3: iris-055  
0: iris-055
```

-N: #nodes

```
$> srun -N 4 -l hostname  
3: iris-058  
2: iris-057  
1: iris-056  
0: iris-055
```

-c: #cpus/task ~#thread/task

```
$> srun -c 4 -l hostname  
0: iris-055
```

```
$> srun -N 2 -n 4 -l hostname
```

```
3: iris-056
```

```
0: iris-055
```

```
1: iris-055
```

```
2: iris-055
```

```
$> srun -N 2 --ntasks-per-node 2 -l hostname
```

```
3: iris-056
```

```
2: iris-056
```

```
1: iris-055
```

```
0: iris-055
```

Job submission with salloc

```
$> salloc [-N #node] [-n #task] [--ntasks-per-node #n]
```

```
$> salloc -N 4
salloc: Granted job allocation 59955
salloc: Waiting for resource configuration
salloc: Nodes iris-[055,060-062] are ready for job
$> env | grep SLURM
$> hostname
access1.iris-cluster.uni.lux
$> srun -l hostname
0: iris-055
2: iris-061
1: iris-060
3: iris-062
```



Reservations and scontrol features

```
$> scontrol show job <JOBID>
```

Job info

Reservations and scontrol features

```
$> scontrol show job <JOBID>
```

Job info

```
$> scontrol show {partition,topology}
```

Reservations and scontrol features

```
$> scontrol show job <JOBID>
```

Job info

```
$> scontrol show {partition, topology}
```

```
$> scontrol show reservations
```

Show existing reservations

```
$> scontrol create reservation
```

```
ReservationName=<name>
```

```
accounts=<account_list>
```

```
corecnt=<num>
```

```
duration=[days-]hours:minutes:seconds
```

```
endtime=yyyy-mm-dd[thh:mm[:ss]]
```

```
features=<feature_list>
```

```
flags=maint, overlap, ignore_jobs, daily, weekly
```

```
licenses=<license>
```

```
nodecnt=<count>
```

```
nodes=<node_list>
```

```
partitionname=<partition(s)>
```

```
starttime=yyyy-mm-dd[thh:mm[:ss]]
```

```
users=<user_list>
```


Basic Slurm Launcher Examples

Documentation

https://hpc.uni.lu/users/docs/slurm_launchers.html

See also **PS1**, **PS2** and **PS3**

```
#!/bin/bash -l  
# Request one core for 5 minutes in the batch queue  
  
#SBATCH -N 1  
#SBATCH --ntasks-per-node=1  
#SBATCH --time=0-00:05:00  
#SBATCH -p batch  
#SBATCH --qos=qos-batch
```

[...]

Basic Slurm Launcher Examples (cont.)

```
#!/bin/bash -l
# Request two cores on each of two nodes for 3 hours

#SBATCH -N 2
#SBATCH --ntasks-per-node=2
#SBATCH --time=0-03:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

echo "== Starting run at $(date)"
echo "== Job ID: ${SLURM_JOBID}"
echo "== Node list: ${SLURM_NODELIST}"
echo "== Submit dir. : ${SLURM_SUBMIT_DIR}"

[...]
```

Basic Slurm Launcher Examples (cont.)

```
#!/bin/bash -l
# Request one core and half the memory available on an iris cluster
# node for one day
#
#SBATCH -J MyLargeMemorySequentialJob
#SBATCH --mail-type=end,fail
#SBATCH --mail-user=Your.Email@Address.lu
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH --mem=64GB
#SBATCH --time=1-00:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

echo "== Starting run at $(date)"
echo "== Job ID: ${SLURM_JOBID}"
echo "== Node list: ${SLURM_NODELIST}"
echo "== Submit dir. : ${SLURM_SUBMIT_DIR}"
```

pthread/OpenMP Slurm Launcher

```
#!/bin/bash -l
# Single node, threaded (pthread/OpenMP) application launcher,
# using all 28 cores of an iris cluster node:

#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -c 28
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}
/path/to/your/threaded.app
```

MATLAB Slurm Launcher

```
#!/bin/bash -l
# Single node, multi-core parallel application (MATLAB, Python, R...)
# launcher, using all 28 cores of an iris cluster node:

#SBATCH -N 1
#SBATCH --ntasks-per-node=28
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

module load base/MATLAB
matlab -nodisplay -nosplash < /path/to/inputfile > /path/to/outputfile
```

Intel MPI Slurm Launchers

- Official SLURM guide for Intel MPI

```
#!/bin/bash -l
# Multi-node parallel application IntelMPI launcher,
# using 128 distributed cores:

#SBATCH -n 128
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

module load toolchain/intel
export I_MPI_PMI_LIBRARY=/usr/lib64/libpmi.so
srun -n $SLURM_NTASKS /path/to/your/intel-toolchain-compiled-application
```

OpenMPI Slurm Launchers

- Official SLURM guide for Open MPI

```
#!/bin/bash -l
# Multi-node parallel application openMPI launcher,
# using 128 distributed cores:

#SBATCH -n 128
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

module load toolchain/foss
mpirun -n $SLURM_NTASKS /path/to/your/foss-toolchain-compiled-application
```


Hybrid IntelMPI+OpenMP Launcher

```
#!/bin/bash -l
# Multi-node hybrid application IntelMPI+OpenMP launcher,
# using 28 threads per node on 10 nodes (280 cores):

#SBATCH -N 10
#SBATCH --ntasks-per-node=1
#SBATCH -c 28
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

module load toolchain/intel
export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}
export I_MPI_PMI_LIBRARY=/usr/lib64/libpmi.so
srun -n $SLURM_NTASKS /path/to/your/parallel-hybrid-app
```


Typical Workflow on UL HPC resources



• Preliminary setup

① Connect to the frontend

② Synchronize you code

③ Reserve a few interactive resources

or,


✓ (eventually) build your program

✓ Test on small size problem

✓ Prepare a launcher script

ssh, screen
scp/rsync/svn/git
oarsub -I [...]
on iris: srun -p interactive [...]
gcc/icc/mpicc/nvcc..
mpirun/srun/python/sh..
<launcher>.{sh|py}

Typical Workflow on UL HPC resources



• Preliminary setup

① Connect to the frontend

② Synchronize you code

③ Reserve a few interactive resources

or,

✓ (eventually) build your program

✓ Test on small size problem

✓ Prepare a launcher script

ssh, screen

scp/rsync/svn/git

oarsub -I [...]

on iris: srun -p interactive [...]

gcc/icc/mpicc/nvcc..

mpirun/srun/python/sh...

<launcher>.{sh|py}

• Real Experiment

① Reserve passive resources

oarsub [...] <launcher>

or,

on iris: sbatch -p {batch|long} [...] <launcher>

② Grab the results

scp/rsync/svn/git ..