

allinea



Leaders in parallel software development tools

Allinea Unified environment

More time computing, less time in tools



UNIVERSITÉ DU
LUXEMBOURG

www.allinea.com

Agenda



14:30 – 14:45 : Introduction to Allinea tools in University of Luxembourg

14:45 – 15:45 : Getting started with Allinea DDT (hands-on)


15:45 – 16:30 : Getting started with Allinea MAP (hands-on)

16:30 – 17:30 : Allinea open discussion and day 2 closing.


Discover the new tool “Allinea Performance Reports” !

Try Allinea tools with your own codes !



And now...



Let's talk about us!



New technologies, more parallelism



Need to dive into the code ?

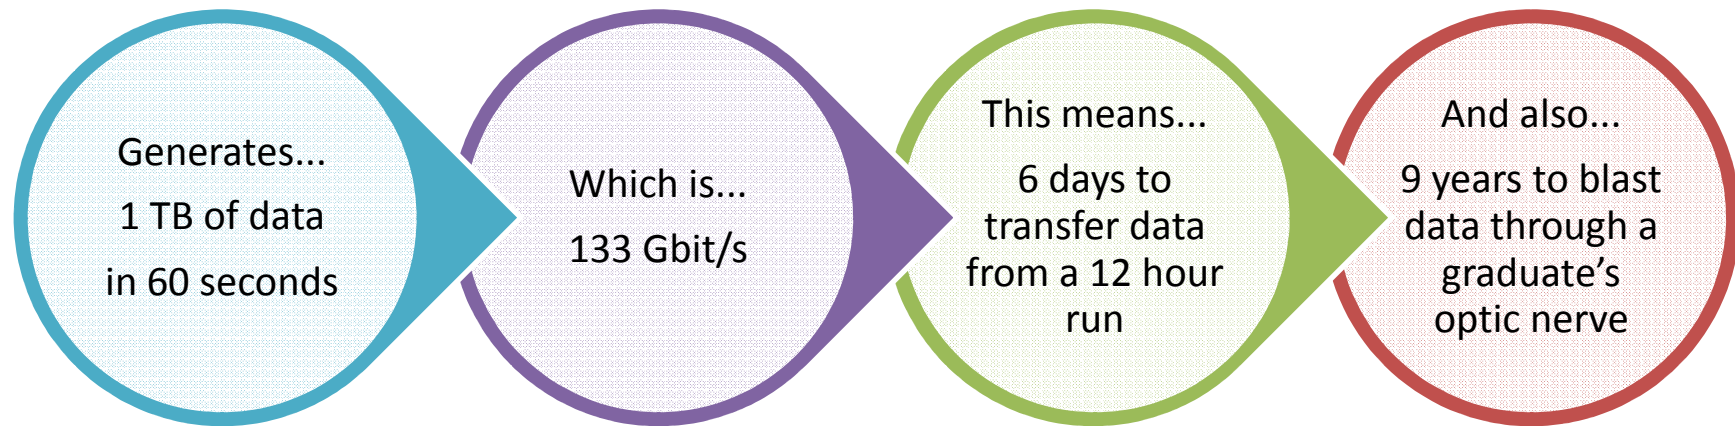
- A modern integrated environment for HPC developers
- Supporting the lifecycle of application development and improvement
 - Allinea DDT : Productively debug code
 - Allinea MAP : Enhance application performance
- Designed for productivity
 - Consistent easy to use tools
 - Fewer failed jobs
- Available at University of Luxembourg
 - Allinea Unified Supercomputing on 64 procs with accelerator support



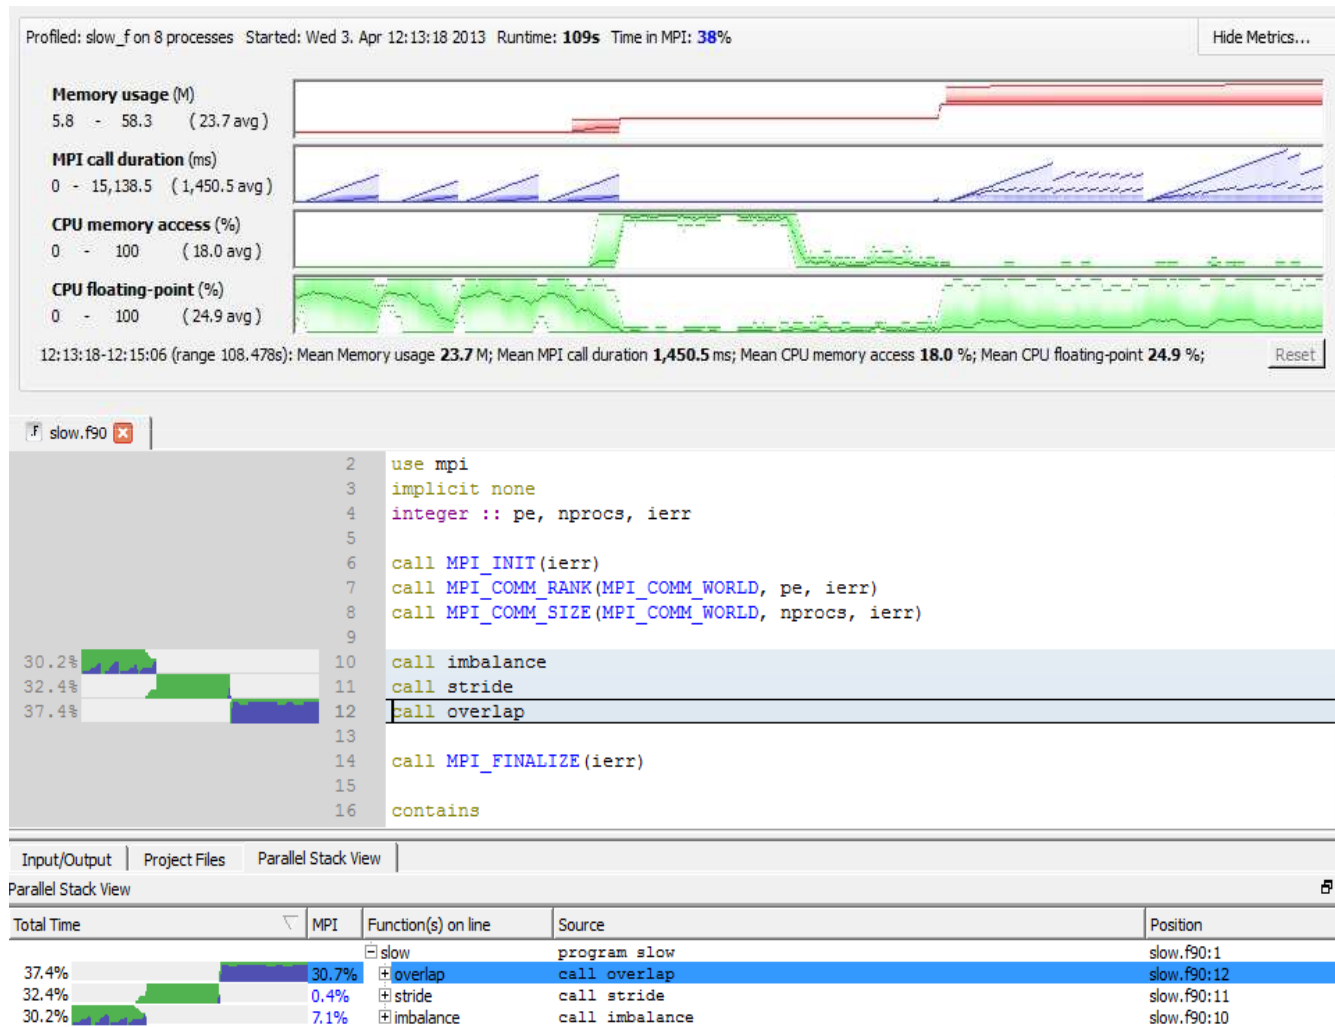
Beware exploding bandwidth needs...



Trivial 16k processes wave equation code running on Titan



Attacking Visual Scalability



Common
horizontal axis



Aggregate across
all processes



Highlight
imbalance visually



Always refer to
source code

Statistic sampling or tracing ?

Complementary approaches

Optimize with
Allinea MAP

- Characterize performance at-scale with a lightweight tool
- See which lines of code are hotspots
- Identify common problems at once

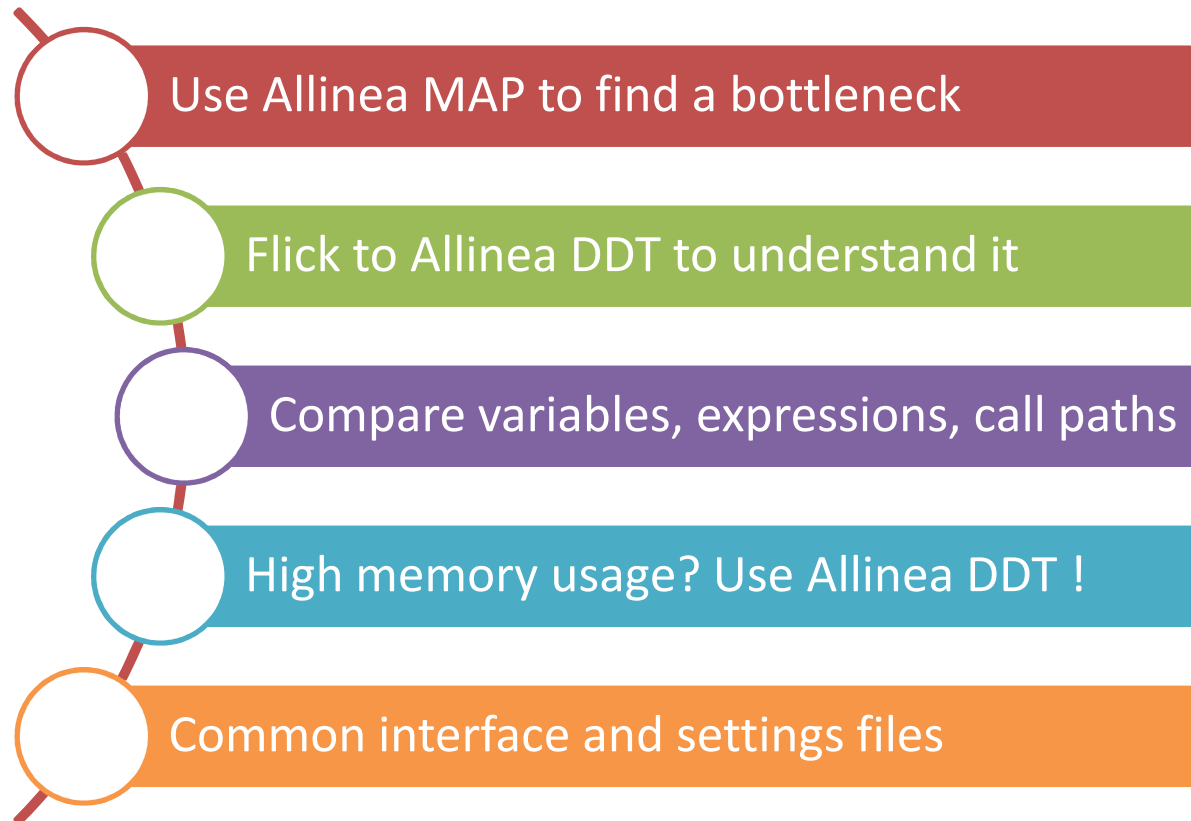
Prepare strategy
with Allinea MAP

- Pass more obscure problems to an expert
- Identify loop(s) to instrument
- Identify performance counter(s) to record

Record traces

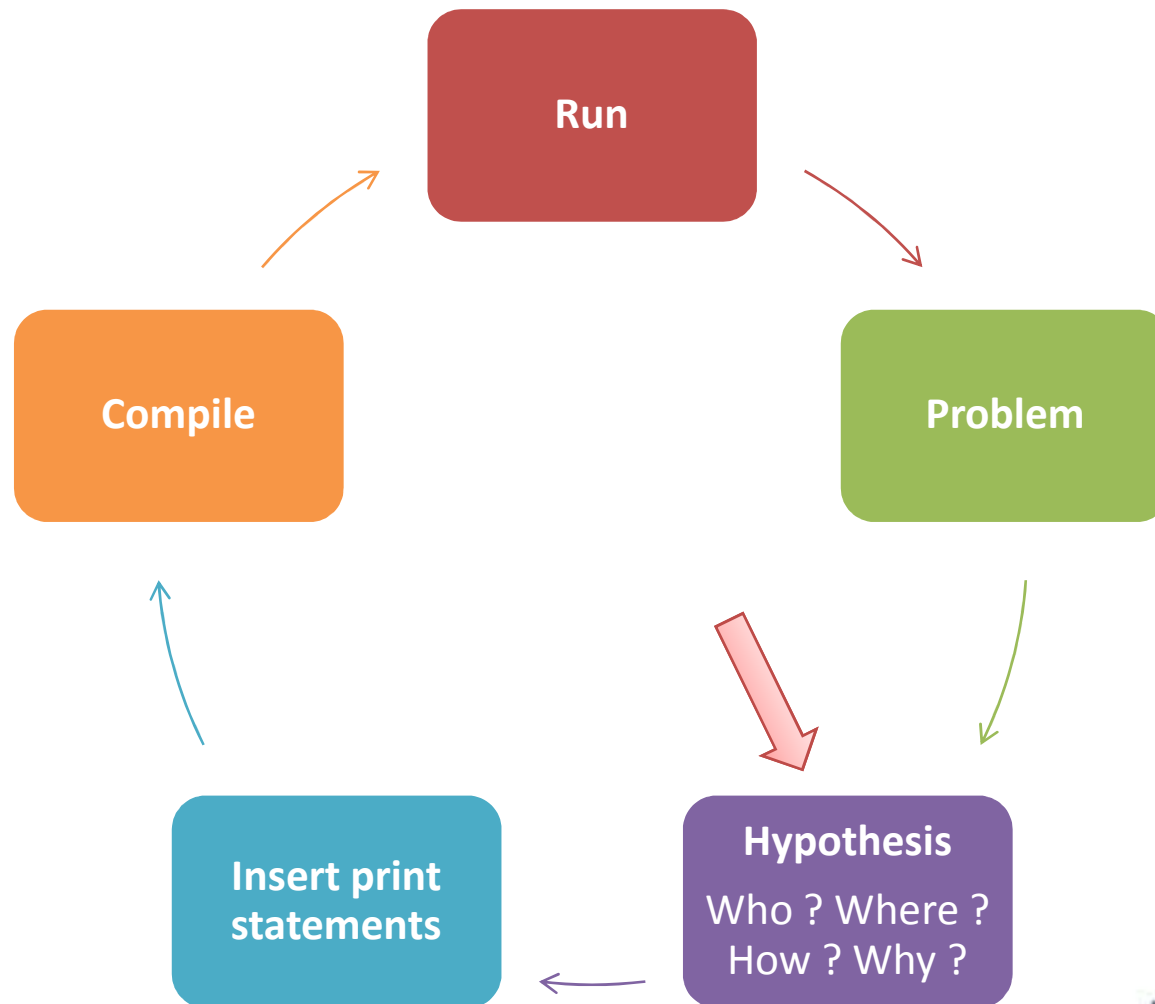
- Retrieve low level details
 - without generating huge traces
 - without huge overheads

Integrated with Allinea DDT



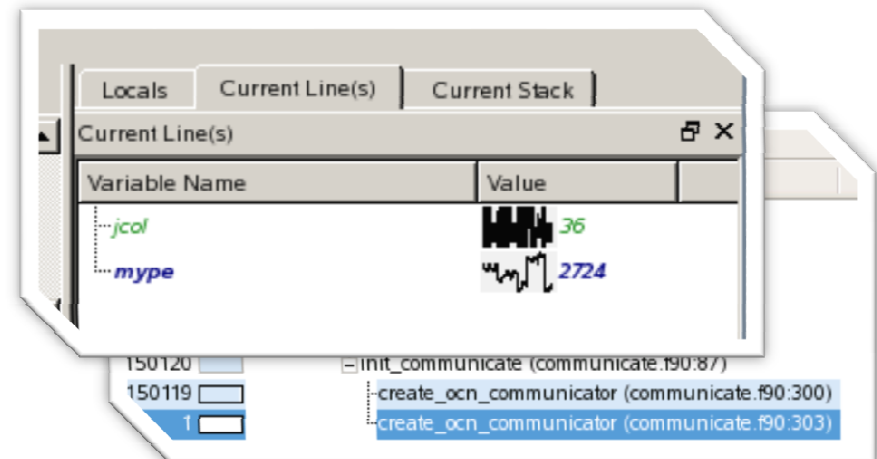
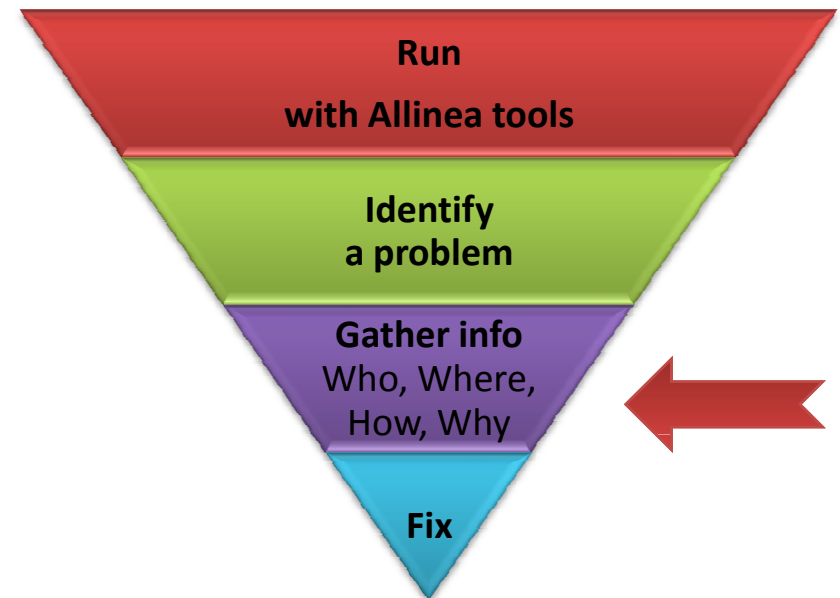
Debugging in practice

The usual method



Allinea DDT helps to understand

- **Who had a rogue behavior ?**
 - Merges stacks from processes and threads
- **Where did it happen?**
 - Allinea DDT leaps to source automatically
- **How did it happen?**
 - Detailed error message given to the user
 - Some faults evident instantly from source
- **Why did it happen?**
 - Unique “Smart Highlighting”
 - Coloring differences and changes
 - Sparklines comparing data across processes



Summary



- **To “make” science quickly, all HPC aspects need to be accessible**
 - Tools need to be usable to avoid wasting time
 - Provided information needs to be adequate
- **Allinea DDT and Allinea MAP : 2 sides of the same coin**
 - Unified profiling and debugging to fix or optimize code
 - Integrates new features to help reduce your time developing code
- **BONUS : Allinea Performance Reports for the HPC users**
 - New product released a few weeks ago – will be available soon at Uni Luxembourg!
 - Understand application behaviour quickly

allinea



Leaders in parallel software development tools

Hands-on workshop

In the beginning was the Word



1- Connect to the front-end node :

```
$ ssh -Y -p 8022 <username>@access-gaia.uni.lu
```

2- Get the training package

```
$ cp /path/to/archive/allinea_wshop.tar.gz $WORK/
```

3- Connect on a compute nodeSetup the environment

```
$ oarsub -I
```

4- Setup the environment

```
$ cd $WORK
```

```
$ tar xvfz $WORK/allinea_wshop.tar.gz
```

```
$ module load OpenMPI/1.6.5-GCC-4.7.2 DDT
```

NOTE : in order to use GCC-4.8.3 (relies on DWARF4), Allinea tools version 4.2 or later is required !

5- Read instructions

```
$ evince allinea_wshop/exercise1/handout_ex1.pdf
```

allinea



Leaders in parallel software development tools

Thank you

Your contacts :

- Technical Support team :
- Sales team :

support@allinea.com

sales@allinea.com

www.allinea.com

Understand cluster usage efficiency

- **Monitors application behavior to provide answers:**

- Are the applications running on the cluster efficient ?
- Are there software or hardware bottlenecks affecting performance ?
- Is the combination of application parameters optimal ?
- What cluster/scale should the user choose for his job ?

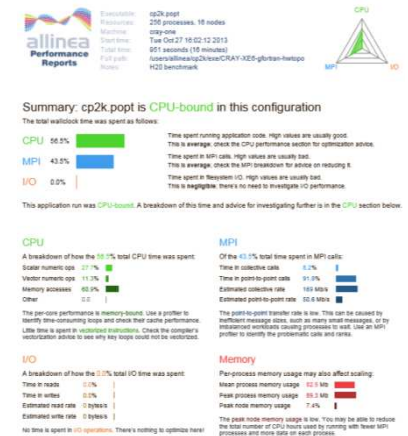
- **Effortless one-touch reports:**

`mpirun -n 42 ./my_executable argument1`


becomes

`perf-report mpirun -n 42 ./my_executable argument1`

- **Fully supported in x86_64 environments**




Better performance quickly and easily



Summary: madbench is I/O-bound in this configuration

The total wallclock time was spent as follows:



This application run was **I/O-bound**. A breakdown of this time and advice for investigating further is in the I/O section below.

CPU

A breakdown of how the 17.9% total CPU time was spent:



The per-core performance is memory-bound. Use a profiler to identify time-consuming loops and check their cache performance.

MPI

Of the 24.5% total time spent in MPI calls:



Most of the time is spent in **collective calls** with a very low transfer rate. This suggests load imbalance is causing synchronization overhead; use an MPI profiler to investigate further.

I/O


A breakdown of how the 47.6% total I/O time was spent:



Most of the time is spent in **write operations** with a low effective transfer rate. This may be caused by contention for the filesystem or inefficient access patterns. Use an I/O profiler to investigate which write calls are affected.

Memory

Per-process memory usage may also affect scaling:



The peak node memory usage is very low. You may be able to reduce the total number of CPU hours used by running with fewer MPI processes and more data on each process.

No instrumentation needed

No need for recompilation or source code

Perfect for ISV applications

Less than 5% runtime overhead

Fully scalable

Run regularly – or in regression tests

Explicit and usable output with hints