

# UL HPC School 2017

## PS2: HPC workflow with sequential jobs (test cases on GROMACS, Java and Python)



---

UL High Performance Computing (HPC) Team

H. Cartiaux

University of Luxembourg (UL), Luxembourg

<http://hpc.uni.lu>

## Latest versions available on **Github**:



UL HPC tutorials:

<https://github.com/ULHPC/tutorials>

UL HPC School:

<http://hpc.uni.lu/hpc-school/>

PS2 tutorial sources:

[http://ulhpc-tutorials.readthedocs.io/en/latest/basic/sequential\\_jobs/](http://ulhpc-tutorials.readthedocs.io/en/latest/basic/sequential_jobs/)





# Summary

- 1 **Introduction**
- 2 Pre-requisites
- 3 Exercise 1: Parametric execution of Gromacs
- 4 Exercise 2: Watermarking images in Python
- 5 Exercise 3: Advanced use case, using a Java program: "JCell"
- 6 Conclusion



# Main Objectives of this Session

- Run sequential, parametric programs on the clusters
- Learn how-to use our set of launcher scripts
- Submit jobs
- use the cluster monitoring tools
  - ↪ Ganglia
  - ↪ Monika & Drawgantt

## Read the full subject of this PS here

- <http://git.io/5cYmPw>



# Summary

- 1 Introduction
- 2 Pre-requisites**
- 3 Exercise 1: Parametric execution of Gromacs
- 4 Exercise 2: Watermarking images in Python
- 5 Exercise 3: Advanced use case, using a Java program: "JCell"
- 6 Conclusion



# Getting started

## 1 Connect to the cluster(s)

```
(laptop)$> ssh {iris,gaia,chaos}-cluster
```

## 2 Send files

```
(laptop)$> rsync -avz local_directory {iris,gaia,chaos}-cluster:
```

## 3 Retrieve files

```
(laptop)$> rsync -avz {iris,gaia,chaos}-cluster:path/to/files  
local_directory
```

## 4 Submit jobs

---

OAR on Chaos/Gaia

Slurm on Iris

---

oarsub -l  
oarsub ./program

srun -p interactive -qos qos-interactive -pty bash  
sbatch program

---



## Tutorial link

**This tutorial is available on github !**

- <https://git.io/vHyh3>



# Summary

- 1 Introduction
- 2 Pre-requisites
- 3 Exercise 1: Parametric execution of Gromacs**
- 4 Exercise 2: Watermarking images in Python
- 5 Exercise 3: Advanced use case, using a Java program: "JCell"
- 6 Conclusion





## Gromacs

### GROMACS: GROningen MACHine for Chemical Simulations


versatile package for molecular dynamics, primarily designed for biochemical molecules

- very large codebase: 1.836.917 SLOC
- many applications in the package, several parallelization modes
- **mdrun**: computational chemistry engine, performing:
  - ↳ molecular dynamics simulations
  - ↳ Brownian Dynamics, Langevin Dynamics
  - ↳ Conjugate Gradient
  - ↳ L-BFGS
  - ↳ Steepest Descents energy minimization
  - ↳ Normal Mode Analysis
- **mdrun** - parallelized using MPI, OpenMP, pthreads and with support for GPU acceleration

## Comparison

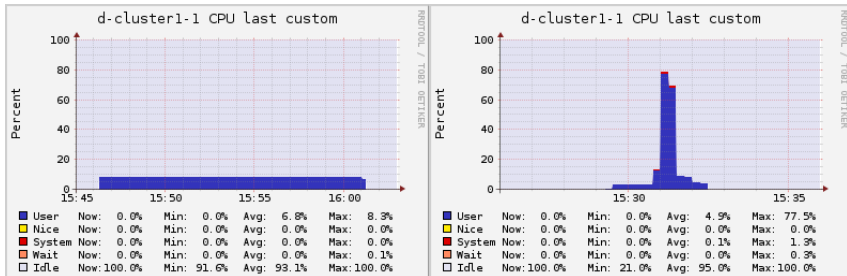
### 2 approaches

- Sequential (loop)
- Parallized (with GNU parallel)





## Comparison - Ganglia






# Summary

- 1 Introduction
- 2 Pre-requisites
- 3 Exercise 1: Parametric execution of Gromacs
- 4 Exercise 2: Watermarking images in Python**
- 5 Exercise 3: Advanced use case, using a Java program: "JCell"
- 6 Conclusion



## Apply a watermark to a given set of pictures

- Simple Python script
- Generic parallel launcher
- Distribute the work on several nodes





# Source image



# Watermarked image





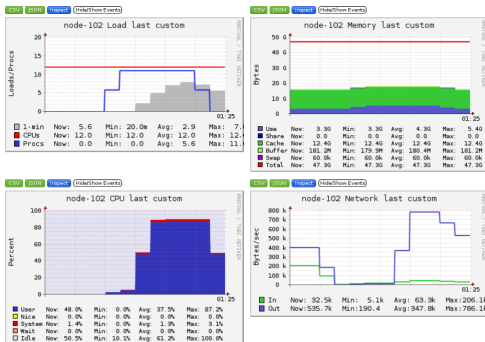
## Summary

- 1 Introduction
- 2 Pre-requisites
- 3 Exercise 1: Parametric execution of Gromacs
- 4 Exercise 2: Watermarking images in Python
- 5 Exercise 3: Advanced use case, using a Java program: "JCell"**
- 6 Conclusion



## Jcell

- **JCell**: a Java framework for working with genetic algorithms
- Example: Generational algorithm for the Combinatorial ECC problem
- Test the variations of these parameters: *Mutation probability* and *Crossover probability*





# Summary

- 1 Introduction
- 2 Pre-requisites
- 3 Exercise 1: Parametric execution of Gromacs
- 4 Exercise 2: Watermarking images in Python
- 5 Exercise 3: Advanced use case, using a Java program: "JCell"
- 6 Conclusion**



## Conclusion

- We have covered one of the most common workflow:  
↳ **parametric jobs**
- Our launchers can be improved!

### Perspectives

- Array jobs
- Best effort jobs
- Checkpoint/Restart mechanism



Thank you for your attention...

# Questions?

<http://hpc.uni.lu>

**The UL High Performance Computing (HPC) Team**  
University of Luxembourg, Belval Campus:  
Maison du Nombre, 4th floor  
2, avenue de l'Université  
L-4365 Esch-sur-Alzette  
*mail:* [hpc@uni.lu](mailto:hpc@uni.lu)



- 1 Introduction
- 2 Pre-requisites
- 3 Exercise 1: Parametric execution of Gromacs
- 4 Exercise 2: Watermarking images in Python
- 5 Exercise 3: Advanced use case, using a Java program: "JCell"
- 6 Conclusion